

# BALANCING REPRESENTATIONS BY IDENTIFYING DEFICITS AND GENERATING EXAMPLES

Anonymous authors

Paper under double-blind review

## ABSTRACT

Self-supervised learning inherits the coverage of its source data. When an unlabeled corpus is long-tailed, the representation becomes dense around frequent modes and thin where coverage is weak. Changing the training objective does not add the missing images. We introduce BRIDGE, Balancing Representations by Identifying Deficits and Generating Examples, a label-free feedback loop around the source set. The current encoder scores every source image by its mean  $k$ NN distance, and a budgeted acquisition rule selects images from a band of that score distribution where support is thin and the generator still produces faithful variants. A frozen image-to-image diffusion model then adds local variants of the selected anchors, and pretraining continues on the enlarged source set. We formulate acquisition as recoverability-weighted allocation of local support and derive when a finite budget should favor a sparse band that excludes the extreme tail and when reallocation after each repair is optimal under a fixed objective. Across five long-tailed, web-scraped and synthetic source regimes, BRIDGE gives the best seven-task linear-probe and  $k$ NN transfer, and paired experiments extend the gain to ViT-S under contrastive and self-distillation objectives, to low-label fine-tuning and to semantic segmentation. The linear-probe gain is smaller under masked reconstruction. Generated repair matches an oracle that restores the withheld real images, and under a frozen encoder the inserted images shrink the neighborhoods of the selected anchors by 4.7 percent against 2.4 percent at controls matched on initial sparsity.




Figure 1: Overview of BRIDGE. Left, an imbalanced source set embedded by the SSL encoder, with sparse under-represented samples in red. Right, the loop. Distance to the  $k$  nearest neighbors in the current latent space selects sparse samples from the upper tail of the score distribution while the extreme tail is excluded, image-to-image generation adds variants of them, and the refined dataset is the source set of the next cycle.

# 1 INTRODUCTION

Self-supervised learning is the standard route to transferable visual representations from unlabeled images (Chen et al., 2020; He et al., 2020; Caron et al., 2021; He et al., 2022; Oquab et al., 2024; Siméoni et al., 2025), and every objective learns from whatever its source corpus contains. When that corpus is long-tailed, the representation concentrates around frequent modes, covers rare regions weakly and transfers worse (Tian et al., 2021; Assran et al., 2023; Kukleva et al., 2023; Jiang et al., 2021). Prior work responds on the model side, with a temperature schedule (Kukleva et al., 2023), a pruned self-competitor (Jiang et al., 2021) or imported external data (Bai et al., 2023), leaving the source set fixed.

We ask whether the source set can be repaired while learning proceeds. BRIDGE, Balancing Representations by Identifying Deficits and Generating Examples, embeds the source set with the current encoder after each cycle and scores every image by its mean distance to its  $k$  nearest neighbors. It selects a sparse but semantically coherent band of that distribution, generates local variants of the selected images with a frozen image-to-image diffusion model and continues pretraining on the enlarged set (Figure 1). The encoder identifies where its own representation is thin, and the generator adds support there. No labels, class frequencies or downstream signals enter the loop.

A finite generation budget must be placed somewhere in the representation, and the placement may or may not be re-estimated after each repair. We formalize both as budgeted allocation over regions of representation space with diminishing returns to local support and repairability that falls in the extreme tail. The optimal band is then sparse and excludes the extreme tail, greedy reallocation is optimal under a fixed objective, and a rule applied once has an effect bounded, qualitatively, by the ratio of budget to existing support. Nine score bands are not separated after one repair, while three acquisition rules separate over five cycles.

Our contributions are as follows.

- A theory of label-free data repair as recoverability-weighted allocation of local support (Section 3), in which sparse regions carry the largest marginal value, greedy reallocation is optimal under a fixed separable objective, and a single-stage allocation has an effect bounded by the ratio of budget to existing support.
- BRIDGE, a label-free feedback loop that instantiates this policy (Section 4) with mean  $k$ NN distance as the support estimate, a trimmed score band that holds sparse images the generator can still vary faithfully, farthest-point sampling against redundancy and a frozen image-to-image generator as the repair operator.
- An evaluation on five source regimes and seven transfer tasks under linear-probe and  $k$ NN protocols (Section 6), with paired ViT-S experiments across four SSL objectives, low-label fine-tuning, semantic segmentation, an oracle arm and matched controls for the closest prior methods.
- An account of where the effect comes from (Sections 6.4 and 6.5). Three acquisition rules are not separated after one cycle of repair and separate after five, and a fixed-encoder intervention test shows the generated support landing at the selected anchors.

## 2 RELATED WORK

**Self-supervised learning under skewed sources.** Contrastive, self-distillation and masked-reconstruction pretraining (Chen et al., 2020; He et al., 2020; Caron et al., 2021; Oquab et al., 2024; Siméoni et al., 2025; He et al., 2022) is sensitive to source imbalance, which has been addressed with clustered experts, the hidden uniform cluster prior and reweighting by estimated class frequency (Tian et al., 2021; Assran et al., 2023; Lin et al., 2023). The closest methods change the objective on a fixed corpus, SDCLR with a pruned self-competitor (Jiang et al., 2021) and Temperature Schedules by emphasizing head and tail at different times (Kukleva et al., 2023), or import external images (Bai et al., 2023).

**Rarity, curation and budgeted acquisition.** Mean deep nearest-neighbor distance has been used as a nonparametric out-of-distribution score (Sun et al., 2022), and we use it as a local-support estimate on the training set itself. Curation by prototype distance or concept-cluster complexity removes

examples (Sorscher et al., 2022; Abbas et al., 2024; Whang et al., 2023), SMOTE interpolates labeled minority examples (Chawla et al., 2002), and DOPING oversamples the low-density edge of a learned latent distribution for anomaly detection with no image generator and no recomputed target (Lim et al., 2018), which our top-tail control reproduces under our generator and budget. Active learning spends a finite budget under a learned representation, its core-set form treating acquisition as  $k$ -center coverage under the greedy farthest-first traversal we use inside the mode window (Sener and Savarese, 2018; Gonzalez, 1985). ProbCover instead maximizes covered probability mass at a fixed radius (Yehuda et al., 2022), and Hacothen et al. (2022) find dense regions the better target at small budgets.

**Generative augmentation and synthetic sources.** SDEdit adds noise to an image and denoises it under a diffusion prior (Ho et al., 2020; Rombach et al., 2022; Meng et al., 2022), and that operator is our repair step. It has been applied to few-shot supervised recognition (Trabucco et al., 2024), class-conditional generation improves supervised ImageNet accuracy (Azizi et al., 2023), and synthetic-data utility depends on prompting, filtering, diversity and domain gap (He et al., 2023; Fan et al., 2024). TADA generates faithful variants of the examples a supervised classifier learns slowly (Nguyen et al., 2026), and we adapt its criterion to a label-free learning-speed signal as a baseline. Synthetic imagery has also been a training source for self-supervised features since Ren and Lee (2018), from procedural noise and image programs (Baradad et al., 2021; 2022) to generator latents, synthetic ImageNet clones and text-to-image views (Jahanian et al., 2022; Sariyildiz et al., 2023; Tian et al., 2023), none allocating a budget over an observed skewed corpus. DiffAug jointly trains an encoder and a conditional diffusion generator so that generated positives improve contrastive learning (Zang et al., 2024), learning that generator from the small skewed source it is meant to repair, while BRIDGE holds a pretrained generator fixed and asks where it should intervene. Industrial data engines close a loop between a deployed model and its training data, diagnosing failures with vision-language and language models and then curating, auto-labeling and verifying (Liang et al., 2024; Worker and Pathak, 2022). They work from supervised task failures with annotation in the loop and acquire by semantic-cluster occupancy, which our inverse cluster-frequency control reproduces. A frozen generator also carries knowledge from its pretraining data (Li et al., 2023; Tang et al., 2023; Yu et al., 2025), so we report a control that distills its features. Appendix A tabulates the closest methods.

### 3 THEORETICAL FRAMEWORK

Targeted repair is modeled as budgeted allocation over coherent regions of representation space. A repair helps only where it adds support in a region carrying downstream risk and the generator preserves the content of its anchor, so the policy weighs support deficit against repairability and redundancy. Let  $r_k(z; Z)$  be the distance from  $z$  to its  $k$ th nearest neighbor in  $Z$  and  $A$  a finite set of valid variants generated from  $z$ . Appendix C gives complete proofs.

**Proposition 1** (Monotone local support). *For any finite  $A$ ,  $r_k(z; Z \cup A) \leq r_k(z; Z)$ . The inequality is strict when an added point enters the first  $k$  neighbors of  $z$  and displaces a neighbor at strictly larger distance.*

Adding points can only introduce smaller candidate distances, so the  $k$ th order statistic cannot increase.

Partition the representation into regions  $j$  with effective support  $n_j$  and model the regional contribution to downstream error by

$$\mathcal{B}(\mathbf{m}) = \sum_j p_j a_j (n_j + m_j + \kappa)^{-\gamma}, \quad (1)$$

where  $p_j$  is the transfer-relevant mass of region  $j$ ,  $a_j$  its difficulty,  $m_j$  the number of valid repairs allocated to it, and  $\kappa, \gamma > 0$ . Write  $\Delta_j(m_j) = p_j a_j [(n_j + m_j + \kappa)^{-\gamma} - (n_j + m_j + 1 + \kappa)^{-\gamma}]$  for the reduction from one more valid repair. None of  $p_j$ ,  $a_j$ ,  $\kappa$ ,  $\gamma$  or the regions is instantiated or estimated in the experiments, and Equation 1 serves as a qualitative model in the diminishing-returns form of empirical scaling laws (Sorscher et al., 2022).

**Proposition 2** (Sparse-region marginal value).  *$\Delta_j(m_j)$  is positive and strictly decreasing in  $n_j + m_j$ . Between two regions of equal transfer mass and difficulty, the region with less effective support has the larger marginal value.*

The forward difference of the convex decreasing map  $x \mapsto x^{-\gamma}$  is positive and decreasing, which gives the claim.

Sparsity alone is not sufficient, because generation fails on corrupt or incoherent anchors. Let  $d$  be a sparsity score,  $b(d)$  the marginal value of a valid repair at that score,  $\rho(d)$  the probability of a faithful variant and  $\lambda \geq 0$  the cost of an unfaithful one, so the expected repair utility is

$$u(d) = \rho(d) b(d) - (1 - \rho(d)) \lambda. \quad (2)$$

**Proposition 3** (Interior recoverable-support optimum). *Suppose  $b$  and  $\rho$  are positive and continuously differentiable on  $[d_{\min}, d_{\max}]$  with  $b'(d) > 0$  and  $\rho'(d) < 0$ , and define  $g(d) = b'(d)/(b(d) + \lambda) + \rho'(d)/\rho(d)$ . If  $g$  is strictly decreasing with  $g(d_{\min}) > 0$  and  $g(d_{\max}) < 0$ , then  $u$  has a unique maximizer  $d^*$  in the interior, characterized by  $g(d^*) = 0$ .*

The sign of  $u'$  equals the sign of  $g$ , so the unique sign change of  $g$  is the unique maximum.

The monotonicity assumption on  $g$  holds for  $b(d) = b_0 d^\beta$  with  $\beta \in (0, 1]$  and  $\rho(d) = \rho_0 e^{-\mu d}$  (Proposition 6), where  $g$  is strictly decreasing for every  $\lambda \geq 0$  and  $d^* = \beta/\mu$  at  $\lambda = 0$ . The optimum moves outward when support value is closer to linear and inward when fidelity decays faster. BRIDGE uses a trimmed score band and estimates neither parameter. For the repeated case, consider a separable objective  $U(\mathbf{m}) = \sum_j U_j(n_j + m_j)$  with each  $U_j$  increasing and discretely concave and a budget of  $R$  repairs,  $\sum_j m_j = R$ .

**Proposition 4** (Adaptive allocation under diminishing returns). *For integer allocations, repeatedly assigning one repair to the region with the largest current marginal gain maximizes  $U(\mathbf{m})$ . A fixed allocation may become suboptimal once a region’s marginal gain falls below that of another region.*

Because every marginal sequence is nonincreasing, the greedy picks are the largest feasible marginals overall. In BRIDGE the encoder changes after every repair, and the regions and their marginals with it, so the benefit of re-estimation in the non-stationary case is an empirical claim.

**Proposition 5** (Single-stage insensitivity). *Let  $\mathbf{m}$  and  $\mathbf{m}'$  be any two allocations of the same budget of  $R$  repairs under Equation 1, computed from one representation and applied without re-estimation. Then*

$$|\mathcal{B}(\mathbf{m}) - \mathcal{B}(\mathbf{m}')| \leq R \gamma \max_j p_j a_j (n_j + \kappa)^{-\gamma-1}, \quad (3)$$

which tends to zero as  $R/n_{\min} \rightarrow 0$  for the least-supported region  $n_{\min}$ .

Both allocations subtract at most  $R$  marginals from the same baseline, and the mean value theorem bounds every marginal.

Proposition 5 bounds the single-stage gap between any two rules by the ratio of the budget to the support of the least-supported region. The constants are not estimated. The budget is  $R = 2,500$  repairs per repair stage, 500 anchors with five variants each, and the least-supported regions of a 10,000-image source hold far fewer images than that, so the ratio is large and the bound does not constrain the experiments quantitatively. Section 6.4 measures the gap directly and finds it inside the three-seed spread. The five-cycle separation lies outside this model, since re-estimation makes each score depend on earlier allocations.

## 4 BRIDGE: BALANCING REPRESENTATIONS BY IDENTIFYING DEFICITS AND GENERATING EXAMPLES

### 4.1 LOOP AND SCORING

Let  $D^{(0)}$  be an unlabeled source set and  $f_\theta$  an SSL encoder. BRIDGE alternates pretraining and repair over  $C$  cycles, with encoder  $f_{\theta^{(c)}}$ , source set  $D^{(c)}$  and embedding set  $Z^{(c)}$  at cycle  $c$ . The encoder is the only learned component, so the loop applies to any objective that yields embeddings. Every cycle continues from the previous encoder with a fresh optimizer schedule and retains all earlier images.

After training on  $D^{(c)}$  we embed each image,  $z_i^{(c)} = f_{\theta^{(c)}}(x_i)$ , normalize to unit length, and compute

$$d_i^{(c)} = \frac{1}{k} \sum_{z \in \mathcal{N}_k(z_i^{(c)}; Z^{(c)} \setminus \{z_i^{(c)}\})} \|z_i^{(c)} - z\|_2^2, \quad (4)$$

where  $\mathcal{N}_k(z; Z)$  is the set of  $k$  nearest neighbors of  $z$  in  $Z$ , so a large score marks an image in a sparse part of the current representation. Scoring uses this squared form. The geometry diagnostic of Section 6.5 uses the unsquared distance to the 20th neighbor under a separate index, a statistic used to measure insertion, with selection based on Equation 4. We use  $k = 100$ , and  $k$  between 10 and 200 moves the seven-task average by about one point (Appendix Table 20).

#### 4.2 SPARSE AND RECOVERABLE ANCHOR SELECTION

The extreme tail holds artifacts, degenerate crops and images the generator cannot vary faithfully, while uniform sampling spends most of its budget in regions that are already dense. In three steps (Figure 2) the selector restricts attention to the band  $U^{(c)} = \{i \mid q_{0.75}^{(c)} \leq d_i^{(c)} \leq q_{0.99}^{(c)}\}$ , locates the center  $\bar{d}^{(c)}$  of the modal bin of a histogram over that band, forms the pool  $P^{(c)} \subseteq U^{(c)}$  of size  $\alpha B$  that minimizes  $\sum_{i \in P^{(c)}} |d_i^{(c)} - \bar{d}^{(c)}|$ , and selects the anchor set  $S^{(c)}$  of size  $B$  from that pool by greedy farthest-point sampling in embedding space. The histogram uses the automatic bin heuristic of numpy, the larger of the Sturges and Freedman-Diaconis bin counts. In every run the modal bin of the right-skewed band sits at its lower edge, so the candidate pool is the  $4B$  scores just above the 75th percentile and the histogram step changes nothing (Figure 2). We use  $\alpha = 4$  and  $B = 500$ .




Figure 2: Score histogram of one ImageNet-100-LT run at cycle 0 with the 75th and 99th percentiles marked. The green pool holds the  $4B$  scores nearest the modal bin of the band, which lies at its lower edge, the red region is the top-tail control, and the 500 selected anchors are overlaid.

#### 4.3 LOCAL GENERATIVE REPAIR AND CONTINUED PRETRAINING

For each anchor we generate  $G = 5$  variants with a frozen image-to-image diffusion model  $g_\phi$ , conditioning on the anchor image under an empty text prompt so that the repair signal comes from the anchor image alone, and set  $D^{(c+1)} = D^{(c)} \cup \{g_\phi(x_i, \epsilon_j) \mid i \in S^{(c)}, j \in [G]\}$ . The noise  $\epsilon_j$  indexes independent translations of one anchor, so variants keep coarse structure while changing texture, pose or color. The main configuration is Stable Diffusion 3 Medium with 20 denoising steps, strength 0.6 and output resolution equal to the training crop. Five cycles contain four repairs, so a full run adds 10,000 images to a source of about 10,000. Algorithm 1 in the appendix gives the loop.

### 5 EXPERIMENTAL SETUP

**Sources, targets and protocol.** We pretrain on five unlabeled source regimes, the long-tailed ImageNet-100-LT (Liu et al., 2019), CIFAR-10-LT and CIFAR-100-LT (Krizhevsky, 2009), and 10,000-image subsets of the web corpora PASS (Asano et al., 2021) and DiffusionDB (Wang et al., 2023). Transfer is measured on seven benchmarks, CIFAR-10-LT, CIFAR-100-LT, Stanford Cars (Krause et al., 2013), FGVC Aircraft (Maji et al., 2013), Oxford Flowers (Nilsback and Zisserman, 2008), Oxford-IIIT Pets (Parkhi et al., 2012) and ImageNet-100-LT, by linear probe and by  $k$ NN classification on frozen features. Every number is a three-seed mean and standard deviation, and paired comparisons share seeds and checkpoints. ResNet-50 is the default backbone, common to the baselines. The baselines train five cycles of 100 epochs on a source that does not grow, while BRIDGE trains the same epochs with  $k = 100$ ,  $B = 500$  and  $G = 5$  on a source that grows from about 10,000 to about 20,000 images, so it receives about 1.5 times the gradient updates of the baseline in Tables 1 and 2. The cycle ablation, the paired ViT-S experiments and the single-stage controls match updates exactly, and the conventional-augmentation control adds the same number of images. The seed standard deviation of a seven-task average is 0.6 to 1.5 points under five cycles, so differences below about one point are inside seed variation. DINO and MAE follow their published ViT recipes on ViT-S. Appendix D gives implementation details.

Table 1: Linear-probe transfer for five source regimes, three label-derived long-tailed sources and two 10k web corpora, on seven downstream tasks, mean  $\pm$  standard deviation over three seeds. Every method pretrains on the source in the left column under the same schedule. The shaded row is BRIDGE and bold marks the best value in a column, without resolving overlapping standard deviations. Hybrids of BRIDGE with TS and SDCLR are in Appendix Tables 9 and 10.

| Source          | Method        | C10                              | C100                             | Cars                             | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|-----------------|---------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| ImageNet-100-LT | SimCLR        | 73.9 $\pm$ 0.5                   | 46.8 $\pm$ 0.4                   | 13.7 $\pm$ 0.3                   | 19.6 $\pm$ 0.2                   | 35.9 $\pm$ 7.2                   | 38.6 $\pm$ 1.5                   | 50.4 $\pm$ 0.5                   | 39.8 $\pm$ 1.5                   |
|                 | TS            | 72.5 $\pm$ 0.4                   | 44.0 $\pm$ 0.5                   | 12.5 $\pm$ 1.2                   | 15.0 $\pm$ 0.1                   | 35.7 $\pm$ 1.8                   | 40.6 $\pm$ 2.2                   | <b>56.7 <math>\pm</math> 1.2</b> | 39.6 $\pm$ 1.1                   |
|                 | SDCLR         | 41.1 $\pm$ 1.7                   | 18.0 $\pm$ 0.6                   | 1.4 $\pm$ 0.3                    | 2.3 $\pm$ 0.5                    | 10.8 $\pm$ 1.7                   | 10.4 $\pm$ 0.3                   | 21.9 $\pm$ 0.6                   | 15.1 $\pm$ 0.8                   |
|                 | BRIDGE (ours) | <b>75.8 <math>\pm</math> 0.2</b> | <b>49.3 <math>\pm</math> 0.2</b> | <b>16.4 <math>\pm</math> 1.2</b> | <b>22.4 <math>\pm</math> 0.4</b> | <b>36.8 <math>\pm</math> 1.9</b> | <b>42.9 <math>\pm</math> 1.0</b> | 53.8 $\pm$ 0.1                   | <b>42.5 <math>\pm</math> 0.7</b> |
| CIFAR-10-LT     | SimCLR        | 67.8 $\pm$ 0.2                   | 38.0 $\pm$ 1.0                   | 6.5 $\pm$ 0.4                    | 16.0 $\pm$ 0.9                   | <b>19.2 <math>\pm</math> 1.6</b> | 24.0 $\pm$ 1.3                   | 32.0 $\pm$ 0.7                   | 29.1 $\pm$ 0.9                   |
|                 | TS            | 68.5 $\pm$ 1.3                   | 36.0 $\pm$ 1.7                   | 4.3 $\pm$ 0.2                    | 10.3 $\pm$ 0.1                   | 13.1 $\pm$ 2.0                   | 19.7 $\pm$ 2.3                   | 31.5 $\pm$ 1.2                   | 26.2 $\pm$ 1.3                   |
|                 | SDCLR         | 39.6 $\pm$ 1.1                   | 15.6 $\pm$ 0.2                   | 1.1 $\pm$ 0.1                    | 2.4 $\pm$ 0.9                    | 12.0 $\pm$ 1.6                   | 9.1 $\pm$ 1.8                    | 16.5 $\pm$ 1.6                   | 13.8 $\pm$ 1.1                   |
|                 | BRIDGE (ours) | <b>78.1 <math>\pm</math> 0.3</b> | <b>49.8 <math>\pm</math> 0.4</b> | <b>9.2 <math>\pm</math> 0.8</b>  | <b>17.7 <math>\pm</math> 1.0</b> | <b>18.2 <math>\pm</math> 1.1</b> | <b>27.4 <math>\pm</math> 2.7</b> | <b>39.0 <math>\pm</math> 0.6</b> | <b>34.2 <math>\pm</math> 1.0</b> |
| CIFAR-100-LT    | SimCLR        | 59.0 $\pm$ 1.5                   | 31.1 $\pm$ 1.6                   | 4.1 $\pm$ 0.3                    | 11.2 $\pm$ 1.0                   | 12.6 $\pm$ 3.4                   | 20.1 $\pm$ 1.0                   | 28.7 $\pm$ 1.3                   | 23.8 $\pm$ 1.4                   |
|                 | TS            | 56.3 $\pm$ 0.2                   | 28.0 $\pm$ 0.1                   | 2.8 $\pm$ 0.6                    | 8.5 $\pm$ 0.2                    | 15.5 $\pm$ 2.5                   | 16.2 $\pm$ 3.1                   | 26.2 $\pm$ 1.8                   | 21.9 $\pm$ 1.2                   |
|                 | SDCLR         | 37.1 $\pm$ 0.8                   | 15.7 $\pm$ 0.2                   | 1.9 $\pm$ 0.1                    | 2.4 $\pm$ 0.7                    | 9.4 $\pm$ 1.0                    | 9.9 $\pm$ 1.7                    | 16.2 $\pm$ 0.3                   | 13.2 $\pm$ 0.7                   |
|                 | BRIDGE (ours) | <b>74.2 <math>\pm</math> 0.5</b> | <b>48.1 <math>\pm</math> 0.4</b> | <b>8.3 <math>\pm</math> 1.0</b>  | <b>18.2 <math>\pm</math> 0.7</b> | <b>24.4 <math>\pm</math> 2.3</b> | <b>27.6 <math>\pm</math> 0.7</b> | <b>39.8 <math>\pm</math> 1.1</b> | <b>34.4 <math>\pm</math> 1.0</b> |
| PASS-10k        | SimCLR        | 72.3 $\pm$ 0.2                   | 45.5 $\pm$ 0.5                   | 11.0 $\pm$ 0.2                   | 17.5 $\pm$ 1.4                   | 33.3 $\pm$ 1.3                   | 33.3 $\pm$ 0.9                   | 44.7 $\pm$ 0.9                   | 36.8 $\pm$ 0.8                   |
|                 | TS            | 71.5 $\pm$ 0.2                   | 43.5 $\pm$ 0.4                   | 10.2 $\pm$ 1.2                   | 17.2 $\pm$ 0.2                   | 33.1 $\pm$ 2.1                   | 36.4 $\pm$ 1.8                   | <b>49.4 <math>\pm</math> 0.4</b> | 37.3 $\pm$ 0.9                   |
|                 | SDCLR         | 42.2 $\pm$ 1.1                   | 17.7 $\pm$ 0.8                   | 1.8 $\pm$ 0.3                    | 2.2 $\pm$ 0.8                    | 10.6 $\pm$ 1.3                   | 8.2 $\pm$ 1.6                    | 20.8 $\pm$ 0.4                   | 14.8 $\pm$ 0.9                   |
|                 | BRIDGE (ours) | <b>73.9 <math>\pm</math> 0.3</b> | <b>47.1 <math>\pm</math> 0.4</b> | <b>13.6 <math>\pm</math> 0.7</b> | <b>20.5 <math>\pm</math> 0.7</b> | <b>37.7 <math>\pm</math> 4.0</b> | <b>37.1 <math>\pm</math> 0.7</b> | 49.1 $\pm$ 0.5                   | <b>39.9 <math>\pm</math> 1.0</b> |
| DiffusionDB-10k | SimCLR        | 74.4 $\pm$ 0.5                   | 48.6 $\pm$ 0.2                   | 10.9 $\pm$ 0.7                   | 17.6 $\pm$ 0.8                   | 31.2 $\pm$ 2.4                   | 36.1 $\pm$ 1.3                   | 45.3 $\pm$ 0.5                   | 37.7 $\pm$ 0.9                   |
|                 | TS            | 72.3 $\pm$ 0.3                   | 44.4 $\pm$ 0.3                   | 10.7 $\pm$ 0.6                   | 15.4 $\pm$ 0.6                   | 33.1 $\pm$ 1.1                   | <b>39.9 <math>\pm</math> 2.7</b> | 48.8 $\pm$ 0.2                   | 37.8 $\pm$ 0.8                   |
|                 | SDCLR         | 42.4 $\pm$ 0.2                   | 18.4 $\pm$ 0.2                   | 1.1 $\pm$ 0.2                    | 2.4 $\pm$ 0.1                    | 13.0 $\pm$ 3.7                   | 8.2 $\pm$ 1.4                    | 21.0 $\pm$ 0.6                   | 15.2 $\pm$ 0.9                   |
|                 | BRIDGE (ours) | <b>75.9 <math>\pm</math> 0.4</b> | <b>49.9 <math>\pm</math> 0.5</b> | <b>12.8 <math>\pm</math> 1.0</b> | <b>20.0 <math>\pm</math> 0.4</b> | <b>35.7 <math>\pm</math> 1.1</b> | <b>37.7 <math>\pm</math> 0.8</b> | <b>49.1 <math>\pm</math> 0.5</b> | <b>40.1 <math>\pm</math> 0.7</b> |

Table 2:  $k$ NN transfer for the five source regimes and methods of Table 1, classified by  $k$ NN on frozen features without a learned probe. Mean  $\pm$  standard deviation over three seeds, shading and bold as in Table 1.

| Source          | Method        | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|-----------------|---------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| ImageNet-100-LT | SimCLR        | 68.6 $\pm$ 0.4                   | 37.3 $\pm$ 0.2                   | 6.0 $\pm$ 0.7                   | 11.2 $\pm$ 0.4                   | 32.2 $\pm$ 0.7                   | 23.6 $\pm$ 0.3                   | 43.9 $\pm$ 0.9                   | 31.8 $\pm$ 0.5                   |
|                 | TS            | 68.4 $\pm$ 0.3                   | 36.5 $\pm$ 0.1                   | 4.7 $\pm$ 0.2                   | 8.3 $\pm$ 0.5                    | 26.6 $\pm$ 1.0                   | 23.6 $\pm$ 1.7                   | <b>52.2 <math>\pm</math> 0.5</b> | 31.5 $\pm$ 0.6                   |
|                 | SDCLR         | 32.5 $\pm$ 0.4                   | 10.9 $\pm$ 0.2                   | 1.0 $\pm$ 0.3                   | 2.1 $\pm$ 0.1                    | 10.7 $\pm$ 0.7                   | 4.7 $\pm$ 0.2                    | 11.2 $\pm$ 0.2                   | 10.4 $\pm$ 0.3                   |
|                 | BRIDGE (ours) | <b>69.0 <math>\pm</math> 0.5</b> | <b>37.8 <math>\pm</math> 0.3</b> | <b>7.0 <math>\pm</math> 0.2</b> | <b>11.7 <math>\pm</math> 0.2</b> | <b>36.5 <math>\pm</math> 0.4</b> | <b>25.3 <math>\pm</math> 1.0</b> | 48.1 $\pm$ 0.6                   | <b>33.6 <math>\pm</math> 0.4</b> |
| CIFAR-10-LT     | SimCLR        | 66.9 $\pm$ 0.4                   | 33.4 $\pm$ 0.3                   | 3.7 $\pm$ 0.4                   | 9.8 $\pm$ 0.7                    | 16.5 $\pm$ 0.8                   | 13.9 $\pm$ 0.8                   | 24.0 $\pm$ 0.8                   | 24.0 $\pm$ 0.6                   |
|                 | TS            | 66.3 $\pm$ 0.6                   | 27.8 $\pm$ 0.4                   | 2.4 $\pm$ 0.6                   | 6.4 $\pm$ 0.2                    | 10.5 $\pm$ 0.7                   | 11.1 $\pm$ 0.4                   | 22.0 $\pm$ 0.9                   | 20.9 $\pm$ 0.5                   |
|                 | SDCLR         | 30.1 $\pm$ 0.3                   | 10.0 $\pm$ 0.3                   | 1.7 $\pm$ 0.3                   | 2.2 $\pm$ 0.2                    | 10.1 $\pm$ 0.7                   | 4.4 $\pm$ 0.4                    | 9.8 $\pm$ 0.3                    | 9.7 $\pm$ 0.4                    |
|                 | BRIDGE (ours) | <b>74.8 <math>\pm</math> 0.4</b> | <b>41.0 <math>\pm</math> 0.4</b> | <b>4.8 <math>\pm</math> 0.5</b> | <b>10.0 <math>\pm</math> 0.2</b> | <b>24.2 <math>\pm</math> 0.6</b> | <b>16.2 <math>\pm</math> 0.7</b> | <b>29.7 <math>\pm</math> 0.3</b> | <b>28.7 <math>\pm</math> 0.4</b> |
| CIFAR-100-LT    | SimCLR        | 59.3 $\pm$ 0.8                   | 30.1 $\pm$ 0.9                   | 3.0 $\pm$ 0.3                   | 7.4 $\pm$ 0.4                    | 13.7 $\pm$ 0.2                   | 10.8 $\pm$ 0.4                   | 21.8 $\pm$ 1.1                   | 20.9 $\pm$ 0.6                   |
|                 | TS            | 53.1 $\pm$ 0.6                   | 24.6 $\pm$ 0.4                   | 2.4 $\pm$ 0.4                   | 5.5 $\pm$ 0.3                    | 9.3 $\pm$ 0.4                    | 8.3 $\pm$ 0.2                    | 18.7 $\pm$ 1.0                   | 17.4 $\pm$ 0.5                   |
|                 | SDCLR         | 25.7 $\pm$ 1.2                   | 8.5 $\pm$ 0.6                    | 0.9 $\pm$ 0.4                   | 1.9 $\pm$ 0.2                    | 7.3 $\pm$ 0.3                    | 3.9 $\pm$ 0.3                    | 8.3 $\pm$ 1.1                    | 8.1 $\pm$ 0.6                    |
|                 | BRIDGE (ours) | <b>69.3 <math>\pm</math> 0.5</b> | <b>39.5 <math>\pm</math> 0.1</b> | <b>4.8 <math>\pm</math> 0.6</b> | <b>8.9 <math>\pm</math> 0.1</b>  | <b>21.9 <math>\pm</math> 0.1</b> | <b>15.5 <math>\pm</math> 0.1</b> | <b>29.1 <math>\pm</math> 0.7</b> | <b>27.0 <math>\pm</math> 0.3</b> |
| PASS-10k        | SimCLR        | 67.2 $\pm$ 0.2                   | <b>37.0 <math>\pm</math> 0.3</b> | 5.4 $\pm$ 0.3                   | 10.3 $\pm$ 0.5                   | 30.4 $\pm$ 0.0                   | 16.7 $\pm$ 0.5                   | 38.3 $\pm$ 0.3                   | 29.3 $\pm$ 0.3                   |
|                 | TS            | 65.2 $\pm$ 0.7                   | 34.6 $\pm$ 0.3                   | 4.4 $\pm$ 0.2                   | 7.9 $\pm$ 0.2                    | 24.5 $\pm$ 0.1                   | 14.4 $\pm$ 0.4                   | 39.3 $\pm$ 0.4                   | 27.2 $\pm$ 0.3                   |
|                 | SDCLR         | 32.4 $\pm$ 1.5                   | 11.5 $\pm$ 0.8                   | 1.3 $\pm$ 0.2                   | 2.0 $\pm$ 0.1                    | 8.6 $\pm$ 0.7                    | 5.2 $\pm$ 0.2                    | 11.1 $\pm$ 0.8                   | 10.3 $\pm$ 0.6                   |
|                 | BRIDGE (ours) | <b>67.5 <math>\pm</math> 0.3</b> | 36.9 $\pm$ 0.3                   | <b>5.8 <math>\pm</math> 0.5</b> | <b>10.7 <math>\pm</math> 0.4</b> | <b>33.7 <math>\pm</math> 0.2</b> | <b>17.8 <math>\pm</math> 0.6</b> | <b>41.4 <math>\pm</math> 0.8</b> | <b>30.6 <math>\pm</math> 0.4</b> |
| DiffusionDB-10k | SimCLR        | 67.3 $\pm$ 0.1                   | 36.2 $\pm$ 0.4                   | 4.3 $\pm$ 0.2                   | 8.8 $\pm$ 0.1                    | 26.8 $\pm$ 0.4                   | 17.3 $\pm$ 0.1                   | 36.7 $\pm$ 0.8                   | 28.2 $\pm$ 0.3                   |
|                 | TS            | 64.7 $\pm$ 0.4                   | 33.6 $\pm$ 0.3                   | 3.1 $\pm$ 0.5                   | 7.3 $\pm$ 0.3                    | 23.2 $\pm$ 0.4                   | 17.0 $\pm$ 0.0                   | 37.8 $\pm$ 0.6                   | 26.7 $\pm$ 0.4                   |
|                 | SDCLR         | 32.6 $\pm$ 1.2                   | 10.0 $\pm$ 0.3                   | 1.3 $\pm$ 0.1                   | 2.2 $\pm$ 0.2                    | 6.1 $\pm$ 0.7                    | 4.8 $\pm$ 0.4                    | 10.6 $\pm$ 0.5                   | 9.7 $\pm$ 0.5                    |
|                 | BRIDGE (ours) | <b>68.3 <math>\pm</math> 0.2</b> | <b>37.4 <math>\pm</math> 0.4</b> | <b>5.4 <math>\pm</math> 0.3</b> | <b>9.8 <math>\pm</math> 0.1</b>  | <b>30.4 <math>\pm</math> 0.2</b> | <b>18.6 <math>\pm</math> 0.2</b> | <b>39.4 <math>\pm</math> 1.0</b> | <b>29.9 <math>\pm</math> 0.3</b> |

**Compared methods and controls.** The source-regime comparison covers source-matched SimCLR, TS, SDCLR and BRIDGE, with hybrids in the appendix. All ablations use ImageNet-100-LT source pretraining, the same downstream suite and three seeds under the full five-cycle protocol. The generator ablation includes an oracle that restores 2,500 randomly chosen withheld real images per repair, the same count as the generated variants and never a duplicate of an anchor, placing true support at the true deficit. No label-free method has that information, and the comparison asks whether label-free repair does as well at the same budget.

## 6 RESULTS

### 6.1 TRANSFER ACROSS FIVE SOURCE REGIMES

BRIDGE has the best seven-task average in every source regime under both protocols (Tables 1 and 2), and against SimCLR and TS it is best in 31 of 35 linear-probe cells and 33 of 35  $k$ NN cells. The gain over source-matched SimCLR is above ten linear-probe points under CIFAR-100-LT and above five under CIFAR-10-LT, the two smallest and most skewed sources, and it falls between 2.4

Table 3: Paired five-cycle ViT-S experiments on ImageNet-100-LT across four SSL objectives, and end-to-end adaptation of the SimCLR ResNet-50 representations. Both arms of every pair branch from one source checkpoint and receive the same optimizer updates after the branch. Panel A reports per-task linear-probe accuracy with the seven-task linear-probe and  $k$ NN averages, mean  $\pm$  standard deviation over three seeds, and per-task values are in Appendix Tables 24 and 25. Panel B reports full fine-tuning with 1 and 10 percent of the ImageNet-100-LT labels and PASCAL VOC segmentation, each over three paired seeds. Bold marks the better arm of each pair, the shaded row is the arm with BRIDGE, and the two average columns are means of the seven task means without a standard deviation.

| A. Frozen ViT-S transfer under four objectives |              |                                  |                                  |                                 |                                  |                                  |                                  |                                  |             |             |
|------------------------------------------------|--------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|-------------|-------------|
| Objective                                      | Method       | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg. linear | Avg. $k$ NN |
| SimCLR                                         | SSL baseline | 60.5 $\pm$ 0.5                   | 35.0 $\pm$ 0.2                   | 7.2 $\pm$ 0.2                   | 10.9 $\pm$ 0.4                   | 32.9 $\pm$ 2.7                   | 25.0 $\pm$ 1.2                   | 42.0 $\pm$ 1.4                   | 30.5        | 30.4        |
|                                                | with BRIDGE  | <b>60.9 <math>\pm</math> 0.3</b> | <b>36.3 <math>\pm</math> 0.1</b> | <b>8.2 <math>\pm</math> 0.4</b> | <b>11.8 <math>\pm</math> 0.4</b> | <b>36.6 <math>\pm</math> 2.9</b> | <b>29.8 <math>\pm</math> 3.0</b> | <b>43.9 <math>\pm</math> 1.1</b> | <b>32.5</b> | <b>31.1</b> |
| MoCo v3                                        | SSL baseline | 60.0 $\pm$ 0.4                   | 34.5 $\pm$ 0.0                   | 7.6 $\pm$ 0.2                   | 11.1 $\pm$ 1.3                   | 30.3 $\pm$ 1.4                   | 29.3 $\pm$ 3.1                   | 44.8 $\pm$ 1.2                   | 31.1        | 33.0        |
|                                                | with BRIDGE  | <b>62.4 <math>\pm</math> 0.5</b> | <b>36.3 <math>\pm</math> 0.1</b> | <b>9.6 <math>\pm</math> 0.5</b> | <b>13.4 <math>\pm</math> 1.0</b> | <b>36.8 <math>\pm</math> 6.1</b> | <b>31.2 <math>\pm</math> 4.0</b> | <b>47.5 <math>\pm</math> 0.9</b> | <b>33.9</b> | <b>32.8</b> |
| DINO                                           | SSL baseline | 58.9 $\pm$ 0.9                   | 33.6 $\pm$ 1.8                   | 8.1 $\pm$ 0.2                   | 11.4 $\pm$ 1.0                   | <b>38.1 <math>\pm</math> 5.2</b> | 27.5 $\pm$ 2.3                   | 44.6 $\pm$ 1.5                   | 31.8        | 30.5        |
|                                                | with BRIDGE  | <b>61.3 <math>\pm</math> 0.1</b> | <b>35.2 <math>\pm</math> 0.3</b> | <b>9.5 <math>\pm</math> 0.9</b> | <b>13.3 <math>\pm</math> 0.6</b> | 37.2 $\pm$ 2.6                   | <b>34.6 <math>\pm</math> 3.6</b> | <b>46.8 <math>\pm</math> 0.8</b> | <b>34.0</b> | <b>32.0</b> |
| MAE                                            | SSL baseline | 37.7 $\pm$ 3.7                   | 13.2 $\pm$ 1.9                   | 1.5 $\pm$ 0.2                   | 2.4 $\pm$ 1.0                    | <b>10.5 <math>\pm</math> 5.9</b> | <b>10.1 <math>\pm</math> 3.0</b> | 19.9 $\pm$ 2.1                   | 13.6        | 9.6         |
|                                                | with BRIDGE  | <b>41.5 <math>\pm</math> 2.2</b> | <b>15.0 <math>\pm</math> 2.4</b> | <b>1.6 <math>\pm</math> 0.4</b> | <b>3.2 <math>\pm</math> 0.7</b>  | 8.5 $\pm$ 2.3                    | 8.6 $\pm$ 2.2                    | <b>22.2 <math>\pm</math> 2.4</b> | <b>14.4</b> | <b>11.6</b> |

| B. End-to-end downstream adaptation, SimCLR ResNet-50 |                      |                  |                                    |                   |
|-------------------------------------------------------|----------------------|------------------|------------------------------------|-------------------|
| Evaluation                                            | Setting              | SSL baseline     | With BRIDGE                        | Paired difference |
| Full fine-tuning                                      | IN100-LT, 1% labels  | 14.31 $\pm$ 0.33 | <b>16.29 <math>\pm</math> 1.11</b> | +1.97 $\pm$ 0.78  |
| Full fine-tuning                                      | IN100-LT, 10% labels | 21.56 $\pm$ 0.75 | <b>24.72 <math>\pm</math> 1.75</b> | +3.17 $\pm$ 1.66  |
| Segmentation                                          | PASCAL VOC, mIoU     | 12.78 $\pm$ 0.79 | <b>15.20 <math>\pm</math> 0.86</b> | +2.42 $\pm$ 0.27  |

and 3.1 points under the three larger sources, with the CIFAR sources leading under  $k$ NN as well. The seven-task average includes the in-domain target for the three label-derived sources, and on the four fine-grained tasks that are never a source the four-task mean gain over SimCLR is 1.7 to 7.6 linear-probe points. TS wins the in-domain ImageNet-100-LT target under an ImageNet-100-LT or PASS source and Pets under a DiffusionDB source, SimCLR keeps Flowers under CIFAR-10-LT, and TS stays below SimCLR on Cars and Aircraft in every regime. SDCLR’s pruning schedule was tuned for larger sources, and at 10,000 images its seven-task averages fall to 13 to 15. The DiffusionDB source is itself synthetic and carries a generator’s stylistic bias, so the gain there does not come from added natural-image diversity.

## 6.2 BEYOND THE DEFAULT PROTOCOL

Four objectives run paired on ViT-S under their reference recipes and one repair budget, and BRIDGE raises the seven-task linear-probe average of each, by 2.0 points for SimCLR, 2.8 for MoCo v3 and 2.2 for DINO, and by 0.8 for MAE, which is inside seed variation. Under  $k$ NN the MoCo v3 average is unchanged to slightly lower, 33.0 against 32.8, while the MAE average rises 2.0 points and MAE loses on Flowers and Pets under the linear probe. Of the 14 linear-probe and  $k$ NN measurements per objective, 12 improve for SimCLR, MoCo v3 and MAE and 13 for DINO, and the fine-grained targets gain most, up to 7.1 points on Pets for DINO. The same ResNet-50 representations improve after end-to-end adaptation, by 2.0 and 3.2 points of fine-tuning accuracy with 1 and 10 percent of the labels and by 2.4 mIoU on PASCAL VOC segmentation.

## 6.3 ABLATIONS

**Generator.** Generated repair with Stable Diffusion 3 matches the oracle of Section 5 within three-seed variation, 41.5 against 40.7. FLUX at its original sampling settings reaches 35.0. Under matched sampling the single-stage four-task difference between the two generators narrows to within seed variation (Appendix Table 7), and the five-cycle gap is not decomposed further.

**Selection.** The three mode windows with lower cutoffs at the 1st, 50th and 75th percentile reach 42.1 to 42.3 and are not separated at three seeds, while uniform acquisition reaches  $35.5 \pm 1.2$  and the top tail  $34.3 \pm 1.2$ . Welch t-tests on three seeds per arm with Holm adjustment place every window ahead of both at adjusted  $p \leq .011$ . The experiment supports the upper trim of Proposition 3 and leaves its lower boundary unresolved, so the evidence is for a sparse band that excludes the extreme tail. The arms vary trimming and farthest-point diversification together, so the 6.6 to 8.0 points cover both. All five arms use the same generator settings and add the same number of images.

Table 4: Ablations on ImageNet-100-LT under the full five-cycle protocol, linear-probe accuracy, mean and standard deviation over three seeds. Each block changes one component, and its shaded row is the default. The selection, cycle and backbone blocks share one default run (42.3) and the generator block’s default is a separate run set (41.5), a gap inside seed variation. Bold marks the best competing setting in a column, excluding the oracle row, which is the oracle of Section 5 restoring the withheld real images. Table 1 reports the same configuration from a separate export, and the two differ by up to 1.2 points on Cars.  $k$ NN counterparts are in Appendix Table 11.

| Group        | Setting                                 | C10               | C100              | Cars              | Aircraft          | Flowers           | Pets              | IN100-LT          | Avg.              |
|--------------|-----------------------------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|-------------------|
| Generator    | Stable Diffusion 3 (default)            | <b>75.0 ± 0.3</b> | <b>48.5 ± 0.2</b> | <b>16.3 ± 0.8</b> | <b>20.7 ± 0.4</b> | <b>35.7 ± 0.9</b> | <b>41.8 ± 0.8</b> | <b>52.4 ± 0.7</b> | <b>41.5 ± 0.6</b> |
|              | FLUX                                    | 70.7 ± 0.4        | 42.1 ± 1.4        | 10.3 ± 1.3        | 17.5 ± 0.9        | 29.4 ± 4.6        | 31.5 ± 2.9        | 43.7 ± 1.4        | 35.0 ± 1.9        |
|              | Oracle (withheld real images restored)  | 74.8 ± 0.4        | 47.8 ± 0.7        | 13.1 ± 0.1        | 20.5 ± 0.5        | 40.3 ± 6.1        | 38.1 ± 1.5        | 50.3 ± 0.4        | 40.7 ± 1.4        |
| Selection    | Mode window, lower cutoff p1            | 75.6 ± 0.1        | 49.2 ± 0.2        | <b>16.9 ± 0.7</b> | 21.9 ± 0.4        | 34.6 ± 2.9        | 43.5 ± 0.7        | 52.8 ± 0.1        | 42.1 ± 0.7        |
|              | Mode window, lower cutoff p50           | 75.5 ± 0.1        | <b>49.3 ± 0.4</b> | 15.2 ± 0.3        | <b>22.4 ± 0.6</b> | 35.5 ± 1.1        | <b>43.8 ± 3.1</b> | 53.4 ± 0.9        | 42.2 ± 0.9        |
|              | Mode window, lower cutoff p75 (default) | <b>76.0 ± 0.3</b> | 49.1 ± 0.2        | 15.2 ± 1.2        | <b>22.4 ± 0.4</b> | <b>36.8 ± 1.9</b> | 42.9 ± 1.0        | <b>53.8 ± 0.1</b> | <b>42.3 ± 0.7</b> |
|              | Top tail                                | 70.8 ± 0.5        | 43.1 ± 1.0        | 10.0 ± 1.7        | 15.9 ± 1.3        | 27.7 ± 1.0        | 30.5 ± 1.2        | 42.2 ± 1.4        | 34.3 ± 1.2        |
| Cycles       | Uniform                                 | 71.5 ± 0.2        | 43.8 ± 0.7        | 9.3 ± 0.9         | 17.5 ± 1.2        | 32.9 ± 0.4        | 31.7 ± 3.9        | 42.0 ± 0.9        | 35.5 ± 1.2        |
|              | 2 cycles                                | 68.9 ± 0.6        | 41.6 ± 1.3        | 10.6 ± 1.2        | 16.4 ± 0.4        | 30.7 ± 4.5        | 33.2 ± 0.9        | 43.3 ± 0.6        | 35.0 ± 1.4        |
|              | 5 cycles (default)                      | <b>76.0 ± 0.3</b> | <b>49.1 ± 0.2</b> | <b>15.2 ± 1.2</b> | <b>22.4 ± 0.4</b> | <b>36.8 ± 1.9</b> | <b>42.9 ± 1.0</b> | <b>53.8 ± 0.1</b> | <b>42.3 ± 0.7</b> |
| Architecture | 10 cycles                               | 70.5 ± 1.2        | 43.8 ± 0.6        | 9.2 ± 0.9         | 17.2 ± 1.6        | 24.8 ± 5.2        | 31.3 ± 1.7        | 41.9 ± 0.6        | 34.1 ± 1.7        |
|              | ResNet-18                               | 68.8 ± 0.1        | 41.7 ± 0.2        | 13.5 ± 0.3        | 19.7 ± 0.6        | 38.7 ± 3.4        | 38.4 ± 2.6        | 50.2 ± 0.2        | 38.7 ± 1.1        |
|              | ResNet-50 (default)                     | <b>76.0 ± 0.3</b> | <b>49.1 ± 0.2</b> | <b>15.2 ± 1.2</b> | <b>22.4 ± 0.4</b> | <b>36.8 ± 1.9</b> | <b>42.9 ± 1.0</b> | <b>53.8 ± 0.1</b> | <b>42.3 ± 0.7</b> |
|              | ViT-S                                   | 62.7 ± 0.4        | 37.6 ± 0.6        | 8.8 ± 0.5         | 12.2 ± 0.2        | 35.1 ± 4.2        | 30.2 ± 1.6        | 42.5 ± 1.2        | 32.7 ± 1.3        |
|              | ViT-B                                   | 66.0 ± 2.3        | 40.6 ± 2.3        | 8.3 ± 0.5         | 14.3 ± 0.8        | <b>38.8 ± 2.6</b> | 31.3 ± 2.1        | 44.5 ± 1.7        | 34.8 ± 1.8        |




Figure 3: Left, nine bands of the score distribution of unequal percentile widths each receive the whole budget after a single stage, and their averages span 1.3 points (Appendix Table 22). The shaded region is one seed standard deviation of a single band’s seven-task mean, 0.66 points on average and between 0.35 and 0.93, so no pair of bands is separated at three seeds. Right, three acquisition rules under five cycles separate by 8.0 points, re-plotting the selection block of Table 4. Single-stage averages lie near 28 and five-cycle averages between 34 and 42.

**Cycles and backbone.** At fixed total repair volume and optimizer budget, five cycles improve the seven-task average over two by 7.3 points and remain 8.2 points ahead of ten. Two cycles place most of the budget before the representation has changed and ten leave too little training between repairs. Within BRIDGE, ResNet-50 is the strongest backbone at this source size, and Table 3 shows the repair effect on transformers.

**Sensitivity of the acquisition rule.** Appendix Table 20 varies the candidate-pool multiplier, upper cutoff, neighborhood size, distance metric and selection strategy one at a time. Cosine and normalized  $L_2$  scoring select the identical anchor set and differ by 0.4 points on the average, which bounds run-to-run spread, and the table spans 1.3 points with no variant separated from the default at three seeds. That table uses the single-stage protocol, so it bounds the sensitivity of the acquisition rule alone.

#### 6.4 SINGLE-STAGE AGAINST FIVE-CYCLE ACQUISITION

Figure 3 spends one identical budget on each of nine score bands, and after a single stage the densest quartile and the extreme tail are not separated at three seeds, while over five cycles the top tail, uniform acquisition and the mode window separate by 8.0 points. The protocols also differ in repair volume and training, 2,500 images against 10,000, so the comparison does not isolate re-estimation as the cause. Under the single-stage protocol the mode window is not the leader, reaching 18.9 against 20.0 for the TADA-style rule and 19.5 for the top tail on the four-task average, whose per-arm seed standard deviation is 0.1 to 1.0, and no pair of rules separates at three seeds (Appendix Table 6). A targeted rule appears to gain value because each repair changes the representation from which the next anchors are chosen, while the single-stage arms decide once from a fixed model. The DiffAug




Figure 4: Fixed-encoder intervention test, one marker per seed. With the cycle-0 encoder frozen and every generated image inserted into the  $k$ NN index, the radius of the treated anchors falls about twice as far as that of matched controls, so the variants land in their own anchors’ neighborhoods. Purity falls about a third as much at the anchors.




Figure 5: Score distribution over the four repair cycles of one ImageNet-100-LT run, with the 75th to 99th percentile band shaded and its modal bin marked. The distribution shifts to larger scores and broadens at every cycle, so the band is re-derived from the current representation.

adaptation learns its generator jointly with the encoder from the 10,000-image source. It reaches 11.5 and falls further with mode-window anchors, below every arm that adds no generated image, since a generator learned from a small skewed source is weak.

## 6.5 REPRESENTATION GEOMETRY

Under the frozen cycle-0 encoder (Figure 4) we ask whether the generated variants land in the neighborhoods of their own anchors. Inserting the generated images reduces the radius of the treated anchors by 4.7 percent against 2.4 percent for controls matched within class on initial radius, with per-seed ratios of 2.4, 2.1 and 1.4 and a paired difference of  $-2.3$  points on average that is negative on every seed. A variant that missed its anchor’s neighborhood would leave the anchor’s radius unchanged, so the test measures locality together with spillover onto untreated images. Purity falls by 0.6 points at the anchors and by 2.0 at the controls, so the variants enter the anchors’ neighborhoods as same-class support and reach the controls as dilution, with Proposition 1 accounting for the sign of the change alone. The diagnostic index uses  $k = 20$ , so this radius differs from the  $k = 100$  selection score. All 500 cycle-0 anchors have selection scores inside the band (Figure 2), Figure 5 shows why the band is re-derived every cycle, and Appendix B gives the per-seed values (Table 8).

## 7 LIMITATIONS

BRIDGE depends on informative encoder neighborhoods in the first cycle and on faithful local generation under a frozen generator that imports the biases of its own pretraining data. Repair fidelity, the  $\rho$  of Section 3, is not measured in pixel or semantic space. The fixed-encoder test shows that variants land near their anchors, and the deficit of the top-tail arm is consistent with low fidelity there and with the tail being off-distribution for the downstream tasks. Later cycles select generated images as anchors, 7 to 13 percent at cycle 1 and 34 to 45 percent at cycle 4 in the runs we inspected, and no generated-image filter is applied. The band, neighborhood size and budget were held at their defaults across all five sources. In our unoptimized setting the loop raises wall-clock time by roughly 30 to 45 percent, almost all of it diffusion sampling at default step counts without caching (Appendix E). A distilled few-step sampler, cached variants or separate generation hardware would remove most of it, and since the generation budget is fixed while pretraining cost grows with the source, the relative overhead falls with scale, where a fixed budget was not tested.

## 8 CONCLUSION

Under diminishing returns to local support and fidelity that decays in the extreme tail, the budget belongs in a sparse band away from that tail, and reallocation after each repair is optimal under a fixed objective. Without labels, BRIDGE improves transfer on five sources and four objectives, matches the oracle of Section 5 and shrinks the neighborhoods of the images it selects, and its acquisition rule has no measurable effect after one repair and an eight-point separation after five.

## REPRODUCIBILITY STATEMENT

Section 4 and Algorithm 1 specify the loop, Section 5 and Appendix D give every operating parameter, schedule, generator setting and data split, and Appendix C contains complete proofs with their assumptions stated in Section 3. Appendix H reports every per-task number behind the main-text averages. An anonymous code release with configuration files for every experiment is linked in Appendix D.

## ETHICS STATEMENT

The work uses public datasets under their release terms and a frozen pretrained image generator, releases no new generator, scraped collection or generated corpus, and involves no human subjects. Generative repair inherits the biases and memorization risks of the generator’s pretraining data, and a deployment would require a task-specific review of fairness, privacy and content safety, as Appendix G discusses.

## AI USE STATEMENT

Generative AI tools were used to assist with editing the manuscript text, with writing and refactoring analysis and plotting code, and with checking bibliographic metadata against publisher records. They were not used to generate research ideas, experimental results or any number reported in this paper. All AI-assisted text and code were reviewed by the authors, who take responsibility for the final content of this work.

## REFERENCES

- Amro Kamal Mohamed Abbas, Evgenia Rusak, Kushal Tirumala, Wieland Brendel, Kamalika Chaudhuri, and Ari S. Morcos. Effective pruning of web-scale datasets based on complexity of concept clusters. In *International Conference on Learning Representations*, 2024.
- Yuki Asano, Christian Rupprecht, Andrew Zisserman, and Andrea Vedaldi. PASS: An ImageNet replacement for self-supervised pretraining without humans. In Joaquin Vanschoren and Serena Yeung, editors, *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021.
- Mahmoud Assran, Randall Balestriero, Quentin Duval, Florian Bordes, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, and Nicolas Ballas. The hidden uniform cluster prior in self-supervised learning. In *The Eleventh International Conference on Learning Representations*, 2023.
- Shekoofeh Azizi, Simon Kornblith, Chitwan Saharia, Mohammad Norouzi, and David J. Fleet. Synthetic data from diffusion models improves ImageNet classification. *Transactions on Machine Learning Research*, 2023.
- Jianhong Bai, Zuozhu Liu, Hualiang Wang, Jin Hao, Yan Feng, Huanpeng Chu, and Haoji Hu. On the effectiveness of out-of-distribution data in self-supervised long-tail learning. In *The Eleventh International Conference on Learning Representations*, 2023.
- Manel Baradad, Jonas Wulff, Tongzhou Wang, Phillip Isola, and Antonio Torralba. Learning to see by looking at noise. In *Advances in Neural Information Processing Systems*, volume 34, pages 2556–2569, 2021.
- Manel Baradad, Chun-Fu (Richard) Chen, Jonas Wulff, Tongzhou Wang, Rogerio Feris, Antonio Torralba, and Phillip Isola. Procedural image programs for representation learning. In *Advances in Neural Information Processing Systems*, 2022.
- Mathilde Caron, Hugo Touvron, Ishan Misra, Hervé Jégou, Julien Mairal, Piotr Bojanowski, and Armand Joulin. Emerging properties in self-supervised vision transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 9650–9660. IEEE, October 2021. doi: 10.1109/ICCV48922.2021.00951.

- Nitesh V. Chawla, Kevin W. Bowyer, Lawrence O. Hall, and W. Philip Kegelmeyer. SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16: 321–357, 2002. doi: 10.1613/jair.953.
- Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In Hal Daumé III and Aarti Singh, editors, *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 1597–1607. PMLR, 13–18 Jul 2020.
- Lijie Fan, Kaifeng Chen, Dilip Krishnan, Dina Katabi, Phillip Isola, and Yonglong Tian. Scaling laws of synthetic images for model training ... for now. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7382–7392, 2024. doi: 10.1109/CVPR52733.2024.00705.
- Teofilo F. Gonzalez. Clustering to minimize the maximum intercluster distance. *Theoretical Computer Science*, 38:293–306, 1985. doi: 10.1016/0304-3975(85)90224-5.
- Guy Hacohen, Avihu Dekel, and Daphna Weinshall. Active learning on a budget: Opposite strategies suit high and low budgets. In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, volume 162 of *Proceedings of Machine Learning Research*, pages 8175–8195, 2022.
- Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. Momentum contrast for unsupervised visual representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 9729–9738. IEEE, June 2020. doi: 10.1109/CVPR42600.2020.00975.
- Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16000–16009. IEEE, June 2022. doi: 10.1109/CVPR52688.2022.01553.
- Ruifei He, Shuyang Sun, Xin Yu, Chuhui Xue, Wenqing Zhang, Philip Torr, Song Bai, and Xiaojuan Qi. Is synthetic data from generative models ready for image recognition? In *International Conference on Learning Representations*, 2023.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In Hugo Larochelle, Marc’Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 6840–6851. Curran Associates, Inc., 2020.
- Ali Jahanian, Xavier Puig, Yonglong Tian, and Phillip Isola. Generative models as a data source for multiview representation learning. In *International Conference on Learning Representations*, 2022.
- Ziyu Jiang, Tianlong Chen, Bobak J. Mortazavi, and Zhangyang Wang. Self-damaging contrastive learning. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 4927–4939. PMLR, 18–24 Jul 2021.
- Jonathan Krause, Michael Stark, Jia Deng, and Li Fei-Fei. 3d object representations for fine-grained categorization. In *2013 IEEE International Conference on Computer Vision Workshops*, pages 554–561. IEEE, 2013. doi: 10.1109/ICCVW.2013.77.
- Alex Krizhevsky. Learning multiple layers of features from Tiny Images. Technical Report TR-2009, University of Toronto, 2009.
- Anna Kukleva, Moritz Böhle, Bernt Schiele, Hilde Kuehne, and Christian Rupprecht. Temperature schedules for self-supervised contrastive methods on long-tail data. In *The Eleventh International Conference on Learning Representations*, 2023. Published as a conference paper at ICLR 2023.

- Daiqing Li, Huan Ling, Amlan Kar, David Acuna, Seung Wook Kim, Karsten Kreis, Antonio Torralba, and Sanja Fidler. DreamTeacher: Pretraining image backbones with deep generative models. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 16698–16708, 2023. doi: 10.1109/ICCV51070.2023.01531.
- Mingfu Liang, Jong-Chyi Su, Samuel Schuster, Sparsh Garg, Shiyu Zhao, Ying Wu, and Manmohan Chandraker. AIDE: An automatic data engine for object detection in autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 14695–14706, 2024. doi: 10.1109/CVPR52733.2024.01392.
- Swee Kiat Lim, Yi Loo, Ngoc-Trung Tran, Ngai-Man Cheung, Gemma Roig, and Yuval Elovici. DOPING: Generative data augmentation for unsupervised anomaly detection with GAN. In *2018 IEEE International Conference on Data Mining (ICDM)*, pages 1122–1127. IEEE, 2018. doi: 10.1109/ICDM.2018.00146.
- Ci-Siang Lin, Min-Hung Chen, and Yu-Chiang Frank Wang. Frequency-aware self-supervised long-tailed learning. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV) Workshops*, pages 963–972. IEEE, October 2023. doi: 10.1109/ICCVW60793.2023.00103.
- Ziwei Liu, Zhongqi Miao, Xiaohang Zhan, Jiayun Wang, Boqing Gong, and Stella X. Yu. Large-scale long-tailed recognition in an open world. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2537–2546. IEEE, June 2019. doi: 10.1109/CVPR.2019.00264.
- Subhansu Maji, Esa Rahtu, Juho Kannala, Matthew B. Blaschko, and Andrea Vedaldi. Fine-grained visual classification of aircraft, 2013.
- Chenlin Meng, Yutong He, Yang Song, Jiaming Song, Jiajun Wu, Jun-Yan Zhu, and Stefano Ermon. SDEdit: Guided image synthesis and editing with stochastic differential equations. In *International Conference on Learning Representations*, 2022.
- Dang Nguyen, Jiping Li, Jinghao Zheng, and Baharan Mirzasoleiman. Do we need all the synthetic data? targeted image augmentation via diffusion models. In *The Fourteenth International Conference on Learning Representations*, 2026. doi: 10.48550/arXiv.2505.21574.
- Maria-Elena Nilsback and Andrew Zisserman. Automated flower classification over a large number of classes. In *2008 Sixth Indian Conference on Computer Vision, Graphics & Image Processing*, pages 722–729. IEEE, 2008. doi: 10.1109/ICVGIP.2008.47.
- Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, Mahmoud Assran, Nicolas Ballas, Wojciech Galuba, Russell Howes, Po-Yao Huang, Shang-Wen Li, Ishan Misra, Michael Rabbat, Vasu Sharma, Gabriel Synnaeve, Hu Xu, Hervé Jégou, Julien Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024. Published 05/2024.
- Omkar M. Parkhi, Andrea Vedaldi, Andrew Zisserman, and C. V. Jawahar. Cats and dogs. In *2012 IEEE Conference on Computer Vision and Pattern Recognition*, pages 3498–3505. IEEE, 2012. doi: 10.1109/CVPR.2012.6248092.
- Zhongzheng Ren and Yong Jae Lee. Cross-domain self-supervised multi-task feature learning using synthetic imagery. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 762–771, 2018. doi: 10.1109/CVPR.2018.00086.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 10684–10695. IEEE, June 2022. doi: 10.1109/CVPR52688.2022.01042.
- Mert Bulent Sariyildiz, Karteek Alahari, Diane Larlus, and Yannis Kalantidis. Fake it till you make it: Learning transferable representations from synthetic ImageNet clones. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8011–8021, 2023. doi: 10.1109/CVPR52729.2023.00774.

- Ozan Sener and Silvio Savarese. Active learning for convolutional neural networks: A core-set approach. In *International Conference on Learning Representations (ICLR)*, 2018.
- Oriane Siméoni, Huy V. Vo, Maximilian Seitzer, Federico Baldassarre, Maxime Oquab, Cijo Jose, Vasil Khalidov, Marc Szafraniec, Seungeun Yi, Michaël Ramamonjisoa, Francisco Massa, Daniel Haziza, Luca Wehrstedt, Jianyuan Wang, Timothée Darcet, Théo Moutakanni, Léo Sentana, Claire Roberts, Andrea Vedaldi, Jamie Tolan, John Brandt, Camille Couprie, Julien Mairal, Hervé Jégou, Patrick Labatut, and Piotr Bojanowski. DINOv3, 2025.
- Ben Sorscher, Robert Geirhos, Shashank Shekhar, Surya Ganguli, and Ari S. Morcos. Beyond neural scaling laws: Beating power law scaling via data pruning. In *Advances in Neural Information Processing Systems*, 2022.
- Yiyu Sun, Yifei Ming, Xiaojin Zhu, and Yixuan Li. Out-of-distribution detection with deep nearest neighbors. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvári, Gang Niu, and Sivan Sabato, editors, *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 20827–20840. PMLR, 17–23 Jul 2022.
- Luming Tang, Menglin Jia, Qianqian Wang, Cheng Perng Phoo, and Bharath Hariharan. Emergent correspondence from image diffusion. In *Advances in Neural Information Processing Systems*, 2023.
- Yonglong Tian, Olivier J. Hénaff, and Aäron van den Oord. Divide and contrast: Self-supervised learning from uncurated data. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 10063–10074. IEEE, October 2021. doi: 10.1109/ICCV48922.2021.00991.
- Yonglong Tian, Lijie Fan, Phillip Isola, Huiwen Chang, and Dilip Krishnan. StableRep: Synthetic images from text-to-image models make strong visual representation learners. In *Advances in Neural Information Processing Systems*, 2023.
- Brandon Trabucco, Kyle Doherty, Max Gurinas, and Ruslan Salakhutdinov. Effective data augmentation with diffusion models. In *The Twelfth International Conference on Learning Representations*, 2024.
- Zijie J. Wang, Evan Montoya, David Munechika, Haoyang Yang, Benjamin Hoover, and Duen Horng Chau. DiffusionDB: A large-scale prompt gallery dataset for text-to-image generative models. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 893–911, Toronto, Canada, July 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.acl-long.51.
- Steven Euijong Whang, Yuji Roh, Hwanjun Song, and Jae-Gil Lee. Data collection and quality challenges in deep learning: A data-centric AI perspective. *The VLDB Journal*, 32(4):791–813, 2023. doi: 10.1007/s00778-022-00775-9.
- Nyla Worker and Bhumin Pathak. Build custom synthetic data generation pipelines with Omniverse Replicator. NVIDIA Technical Blog, 2022. Non-archival systems context.
- Ofer Yehuda, Avi Dekel, Guy Hacohen, and Daphna Weinshall. Active learning through a covering lens. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, 2022.
- Sihyun Yu, Sangkyung Kwak, Huiwon Jang, Jongheon Jeong, Jonathan Huang, Jinwoo Shin, and Saining Xie. Representation alignment for generation: Training diffusion transformers is easier than you think. In *International Conference on Learning Representations*, 2025.
- Zelin Zang, Hao Luo, Kai Wang, Panpan Zhang, Fan Wang, Stan Z. Li, and Yang You. DiffAug: Enhance unsupervised contrastive learning with domain-knowledge-free diffusion-based data augmentation. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 58174–58196. PMLR, 2024.

## A COMPARISON WITH RELATED ACQUISITION METHODS

Classical oversampling and TADA receive labels or supervised learning signals (Table 5). DOPING uses latent rarity, targets anomaly detection and does not close the loop through a changing SSL representation. GenRep and StableRep make synthetic images or latent neighborhoods the source of contrastive views without allocating a repair budget over an observed skewed corpus. DiffAug is the closest iterative SSL antecedent, and it jointly learns a generator for generic positive views without asking where a frozen generator should intervene. AIDE and industrial data-engine systems establish the model-and-data feedback pattern from supervised task failures with VLM and LLM reasoning, annotation and domain-specific verification. BRIDGE studies the label-free case in which local support in the current SSL representation is both the diagnostic and the acquisition signal.

Table 5: Closest methods and the dimensions relevant to the acquisition problem. Feedback means that the learner or task model changes the next acquisition decision, and the methods are cited in Section 2.

| Method               | Labels or task loss         | Acquisition signal                 | Data action                     | Feedback |
|----------------------|-----------------------------|------------------------------------|---------------------------------|----------|
| SMOTE                | class labels                | class minority                     | interpolate examples            | no       |
| DOPING               | no labels                   | latent edge density                | latent oversampling             | no       |
| Concept pruning      | no task labels              | cluster complexity                 | remove examples                 | no       |
| DA-Fusion            | labels or few-shot set      | class and exemplar                 | image-conditioned diffusion     | no       |
| TADA                 | supervised loss dynamics    | slow-learned examples              | faithful diffusion variants     | no       |
| GenRep and StableRep | prompts or generator latent | global sampling                    | synthetic SSL views             | no       |
| DiffAug              | no labels                   | learned positive-view objective    | jointly learned diffusion views | yes      |
| AIDE                 | task labels and failures    | VLM and LLM failure analysis       | curate, generate, annotate      | yes      |
| BRIDGE               | no labels                   | local support in current SSL space | budgeted local variants         | yes      |

Table 6: Acquisition and repair pipelines under one matched single-stage protocol on ImageNet-100-LT. Every arm branches from the same source checkpoint, receives the same optimizer budget and, where it generates, adds the same number of images. Transfer averages linear-probe accuracy on Cars, Aircraft, Flowers and Pets, and 1-NN accuracy is measured on the source set. Mean  $\pm$  standard deviation over three seeds. Four tasks and one stage, so the values are on a different scale from the five-cycle tables (Tables 1 and 4). The shaded row is the default configuration. No value is bolded, since the five acquisition rules lie within 1.7 points, inside single-arm seed variation.

| Family           | Method                                 | Four-task transfer | Source 1-NN      |
|------------------|----------------------------------------|--------------------|------------------|
| Acquisition rule | Uniform acquisition                    | 18.33 $\pm$ 0.32   | 28.97 $\pm$ 2.11 |
|                  | Top tail, DOPING-style                 | 19.51 $\pm$ 0.22   | 28.63 $\pm$ 1.48 |
|                  | Inverse cluster frequency, AIDE-style  | 18.66 $\pm$ 0.10   | 28.66 $\pm$ 2.42 |
|                  | Slowest-learned examples, TADA-style   | 20.00 $\pm$ 0.98   | 28.78 $\pm$ 2.22 |
|                  | Mode window, BRIDGE                    | 18.88 $\pm$ 0.58   | 28.69 $\pm$ 1.82 |
| Repair operator  | Mode window, conventional augmentation | 18.52 $\pm$ 0.65   | 28.91 $\pm$ 1.96 |
|                  | DiffAug adaptation                     | 11.50 $\pm$ 1.87   | 19.39 $\pm$ 2.78 |
|                  | DiffAug adaptation with BRIDGE anchors | 10.42 $\pm$ 2.17   | 20.28 $\pm$ 2.64 |
| Generator prior  | SD3 feature distillation, no images    | 18.02 $\pm$ 0.47   | 26.27 $\pm$ 2.29 |
|                  | SD3 feature distillation with BRIDGE   | 18.87 $\pm$ 0.45   | 29.45 $\pm$ 2.26 |

Table 7: Generator sensitivity on the four-task diagnostic (Cars, Aircraft, Flowers, ImageNet-100-LT) after one repair stage from a common checkpoint, three paired seeds. The original FLUX setting of six steps and guidance 2.5 uses the same generator setting as the 35.0 FLUX run of Table 4. Guidance under an empty prompt gives bit-identical images at 3.0, 5.0 and 7.5.

| Generator      | Setting               | Four-task transfer |
|----------------|-----------------------|--------------------|
| FLUX.1-dev     | 6 steps, guidance 2.5 | 18.23 $\pm$ 0.98   |
| FLUX.1-dev     | 20 steps              | 18.78 $\pm$ 1.43   |
| FLUX.1-schnell | 20 steps              | 18.85 $\pm$ 1.30   |
| FLUX.1-schnell | guidance 3            | 19.68 $\pm$ 0.81   |
| SD3            | 6 steps               | 18.58 $\pm$ 1.14   |
| SD3            | 20 steps (default)    | 19.70 $\pm$ 0.35   |
| SD3            | guidance 1.0 to 7.5   | 18.42 to 18.88     |
| SD3            | strength 0.3 to 0.9   | 18.60 to 19.63     |

Table 8: Fixed-encoder intervention test behind Figure 4. The cycle-0 encoder is frozen, every generated image of the run is inserted into the  $k$ NN index ( $k = 20$ ), and the same original images are queried before and after. Anchors are the images selected at cycle 0, controls are matched one-to-one within class on initial radius and are never selected in any cycle. Radius change is relative to the initial radius, purity is the fraction of the 20 neighbors sharing the query label, and the last two columns are the anchor-minus-control differences.

| Seed | Radius change (%) |          | Purity change (pts) |          | Difference |        |
|------|-------------------|----------|---------------------|----------|------------|--------|
|      | Anchors           | Controls | Anchors             | Controls | Radius     | Purity |
| 0    | -5.38             | -2.27    | -1.57               | -2.79    | -3.11      | +1.22  |
| 1    | -5.01             | -2.35    | -0.20               | -1.91    | -2.66      | +1.70  |
| 2    | -3.62             | -2.58    | -0.07               | -1.17    | -1.05      | +1.09  |
| Mean | -4.67             | -2.40    | -0.61               | -1.96    | -2.27      | +1.34  |




Figure 6: Repair against no repair for four objectives on ViT-S, one bar per downstream measurement (linear probe and  $k$ NN on each of seven datasets), from the paired runs of Table 3. Blue bars are improvements. Red bars are regressions, two for SimCLR and MoCo v3 (CIFAR-10 and CIFAR-100 under  $k$ NN), one for DINO (Flowers) and two for MAE (Flowers and Pets under the linear probe).

## B REPRESENTATION DIAGNOSTICS

The fixed-encoder intervention test freezes the cycle-0 encoder and queries the same original images before and after inserting all generated samples, so any change reflects the inserted points alone. First-cycle anchors are matched one-to-one within class to original images with the closest initial  $k$ NN radius that are never selected in any cycle, and Table 8 reports the per-seed values behind Figure 4. Figure 6 draws every paired ViT-S measurement of Table 3 as a difference. The learned-space comparison retrain the encoder after the intervention and evaluates the same original-image panel. On that panel MoCo v3 on ViT-S gains  $21.97 \pm 1.58$  in effective rank,  $0.296 \pm 0.019$  in spectral entropy and  $3.80 \pm 0.80$  points of 1-NN accuracy after repair, and the first two survive Holm correction, whereas SimCLR improves transfer while its effective rank falls and its local effect is mixed. We therefore treat the fixed-encoder result as evidence about where new support is inserted and the learned-space statistics as objective-dependent diagnostics.

**Where the temperature-schedule hybrid helps.** Combining BRIDGE with TS lowers the seven-task average in every source regime, and it wins 7 of the 70 individual cells (Appendix Tables 9 and 10). Six of those seven fall on CIFAR-10 and ImageNet-100-LT, the two long-tailed downstream targets, and the largest are +3.3 and +2.8 points on ImageNet-100-LT and CIFAR-10 under a PASS and a CIFAR-10-LT source. The fine-grained targets lose in almost every case. TS reweights the contrastive temperature toward rare samples, which sharpens the head-to-tail boundary that the long-tailed targets score, while BRIDGE adds support that broadens transfer. The two interventions serve different downstream targets.

## C PROOFS FOR THE SUPPORT-ALLOCATION FRAMEWORK

**Proposition 6** (Closed-form interior optimum). *Let  $b(d) = b_0 d^\beta$  with  $b_0 > 0$  and  $\beta \in (0, 1]$ , and let  $\rho(d) = \rho_0 e^{-\mu d}$  with  $\mu > 0$ . Then  $g$  is strictly decreasing on  $(0, \infty)$  for every  $\lambda \geq 0$ , so Proposition 3 applies whenever  $g$  changes sign on the interval. If  $\lambda = 0$  the maximizer is  $d^* = \beta/\mu$ , independent of  $b_0$  and  $\rho_0$ .*

*Proof of Proposition 1.* Sort the distances from  $z$  to points in  $Z$  and denote their  $k$ th order statistic by  $r_k(z; Z)$ . Adding  $A$  preserves every old candidate distance and may introduce smaller ones. The  $k$ th order statistic can therefore only decrease. It decreases strictly when an added distance enters the first  $k$  positions and displaces a strictly larger distance.  $\square$

*Proof of Proposition 2.* Let  $x = n_j + m_j + \kappa$ . Since  $f(x) = x^{-\gamma}$  is positive, strictly decreasing and strictly convex for  $x > 0$ , the forward difference  $f(x) - f(x + 1)$  is positive and strictly decreases with  $x$ . Multiplication by  $p_j a_j > 0$  preserves both properties. Equal  $p_j a_j$  therefore leaves effective support as the sole determinant of the marginal ordering.  $\square$

*Proof of Proposition 3.* Rewrite the utility as  $u(d) = \rho(d)(b(d) + \lambda) - \lambda$ . Its derivative is

$$u'(d) = \rho(d)(b(d) + \lambda) \left[ \frac{b'(d)}{b(d) + \lambda} + \frac{\rho'(d)}{\rho(d)} \right] = \rho(d)(b(d) + \lambda) g(d).$$

The factor preceding  $g(d)$  is positive. Since  $g$  is strictly decreasing and changes sign across the interval, it has exactly one root  $d^*$  by continuity. Hence  $u$  increases before  $d^*$  and decreases after it, making  $d^*$  the unique global maximizer.  $\square$

*Proof of Proposition 6.* With  $b(d) = b_0 d^\beta$  and  $\rho(d) = \rho_0 e^{-\mu d}$  we have  $\rho'(d)/\rho(d) = -\mu$ , so

$$g(d) = \frac{b_0 \beta d^{\beta-1}}{b_0 d^\beta + \lambda} - \mu = \frac{\beta d^{\beta-1}}{d^\beta + c} - \mu, \quad c = \lambda/b_0 \geq 0.$$

Write  $h(d) = \beta d^{\beta-1}(d^\beta + c)^{-1}$ . Differentiating,

$$h'(d) = \frac{\beta d^{\beta-2}[(\beta-1)(d^\beta + c) - \beta d^\beta]}{(d^\beta + c)^2} = \frac{\beta d^{\beta-2}[(\beta-1)c - d^\beta]}{(d^\beta + c)^2}.$$

For  $\beta \in (0, 1]$  we have  $(\beta-1)c \leq 0 < d^\beta$  for all  $d > 0$ , so the bracket is strictly negative and  $h'(d) < 0$ . Hence  $g = h - \mu$  is strictly decreasing for every  $\lambda \geq 0$ . When  $\lambda = 0$  we have  $c = 0$  and  $g(d) = \beta/d - \mu$ , whose unique root is  $d^* = \beta/\mu$ . Since  $g$  is positive below it and negative above it,  $d^*$  is the maximizer by the argument of Proposition 3.  $\square$

*Proof of Proposition 4.* For region  $j$  define the marginal sequence  $\delta_{j,t} = U_j(n_j + t) - U_j(n_j + t - 1)$  for  $t \geq 1$ , which discrete concavity makes nonincreasing, so every feasible allocation of  $R$  repairs selects a prefix of length  $m_j$  from each sequence and its utility is the baseline plus the selected prefix marginals. After  $t$  greedy steps the selected set is the  $t$  largest prefix-feasible marginals, and the next greedy pick is the largest remaining prefix-feasible marginal, so the set stays the top  $t + 1$  and the claim follows by induction on  $t$ . No feasible allocation can therefore exceed the greedy total. For the second sentence take two regions with marginal sequences  $(3, 1, 1)$  and  $(2, 2, 2)$  and a budget of three, where the fixed allocation  $(3, 0)$  built from the initial ordering yields 5 while greedy yields  $3 + 2 + 2 = 7$ .  $\square$

*Proof of Proposition 5.* Write  $f(x) = x^{-\gamma}$ . For  $x > 0$  the mean value theorem gives  $f(x) - f(x + 1) = \gamma \xi^{-\gamma-1}$  for some  $\xi \in (x, x + 1)$ , so

$$0 < \Delta_j(m_j) = p_j a_j [f(n_j + m_j + \kappa) - f(n_j + m_j + 1 + \kappa)] \leq p_j a_j \gamma (n_j + \kappa)^{-\gamma-1},$$

using  $\xi > n_j + m_j + \kappa \geq n_j + \kappa$  and  $m_j \geq 0$ . Both  $\mathcal{B}(\mathbf{m})$  and  $\mathcal{B}(\mathbf{m}')$  are obtained from the common baseline  $\mathcal{B}(\mathbf{0})$  by subtracting sums of at most  $R$  such marginals, since  $\sum_j m_j = \sum_j m'_j = R$ . Each sum lies in  $[0, R\gamma \max_j p_j a_j (n_j + \kappa)^{-\gamma-1}]$ , and the difference of two numbers drawn from an interval is bounded by its length. Monotonicity of  $(n + \kappa)^{-\gamma-1}$  in  $n$  gives the  $n_{\min}$  form.  $\square$

## D IMPLEMENTATION AND REPRODUCIBILITY

**Training schedule.** The encoder is warm-started from the previous cycle and is never retrained from scratch. Each cycle uses a fresh optimizer and learning-rate schedule. All earlier real and generated images remain in the source set, so the setting differs from sequential-task learning in which previous data disappear. The cycle ablation holds the total optimizer update budget and the total number of generated images fixed and compares fewer repairs with more learning between them against more frequent repairs with less.

**Core configuration.** The main experiments use ResNet-50 encoders, five cycles of 100 epochs, mean  $k$ NN distance with  $k = 100$ , 500 anchors per cycle and five variants per anchor. The optimizer is AdamW with weight decay  $10^{-4}$ , and method-specific learning rates and temperatures follow the settings of each baseline. The paired ViT-S experiments use the reference recipe of each objective. DINO uses two global and eight local crops at  $96 \times 96$  with a bottleneck projection head, and MAE runs five cycles of 80 epochs.

**Scoring and selection.** Scoring uses the mean squared normalized  $\ell_2$  distance of Equation 4, and raw  $\ell_2$ , cosine and multiscale- $k$  variants are reported as a sensitivity study in Table 20. For unit vectors the squared distance is an affine function of the cosine, which is why the cosine row selects the identical anchor set. The mode-window selector restricts the score distribution to the 75th to 99th percentile range, builds a histogram with the automatic bin heuristic, forms a mode-centered candidate pool of size  $4B$  and applies greedy farthest-point sampling on normalized features to select the final  $B$  anchors.

**Generators.** The main method uses Stable Diffusion 3 Medium in image-to-image mode, where an empty prompt makes the guidance scale inert (Table 7). The generation ablation replaces it with FLUX or with the oracle, which adds the withheld real images and runs no generator. In the PASS and DiffusionDB experiments the SD3 batch size is 12 to fit the 24-hour wall-time limit.

**Single-stage controls.** Every arm of Table 6 branches from one SimCLR ResNet-50 checkpoint per seed, adds 2,500 images where it generates, and trains one stage of at most 100 epochs capped at 4,850 optimizer steps at batch 128. The TADA-style arm ranks source images by the decrease of their SSL loss over the first stage and selects the slowest-learned 500. The inverse cluster-frequency arm clusters the current embeddings and samples anchors with probability inverse to cluster occupancy. The conventional-augmentation arm applies the standard SimCLR augmentation pipeline five times to each mode-window anchor in place of the generator. The DiffAug arm trains a conditional diffusion generator jointly with the encoder on the source, following the published objective, and its BRIDGE variant conditions that generator on the mode-window anchors. The distillation arm adds a feature-matching loss between the encoder and frozen SD3 activations on the source images and adds no image.

**Source data.** PASS and DiffusionDB use `train[:10000]` with pretraining split override `[1.0, 0.0, 0.0]`. The long-tailed ImageNet-100, CIFAR-10 and CIFAR-100 sources use their standard long-tailed configurations without the 10,000-image restriction.

**Seeds and reporting.** All claims rest on three-seed averages reported as mean and standard deviation. Paired comparisons share seeds and, where stated, source checkpoints. Multiple comparisons within one table are corrected by the Holm procedure. No two conditions within a table share an encoder fingerprint.

**Code availability.** An anonymous code release is available at <https://anonymous.4open.science/status/FOMO-0488>.

## E COMPUTE RESOURCES AND RUNTIME

Most experiments use one node with two NVIDIA A100 80GB GPUs under a 24-hour wall-time limit. The DINO arm fits on a single 48GB card at the batch size used. The 454 completed runs retained for this paper total approximately 5,471 wall-clock hours, of which the baselines and controls that appear only in the appendix account for 2,987. Runs on two-GPU nodes consume roughly twice that in GPU hours, while the paired ViT-S runs hold a single card each. Failed runs, timeouts, cache migration and pilot jobs are excluded from these totals.

Mean wall-clock times over three seeds are 16.9 hours for SimCLR and 21.8 for BRIDGE on ImageNet-100-LT, 10.2 and 14.9 on PASS, and 11.6 and 16.4 on DiffusionDB. The loop therefore raises wall-clock time by roughly 30 to 45 percent over the corresponding no-repair run, of which the embedding pass is a small fraction and diffusion sampling the remainder. Four repairs add 10,000

images, which is 10 percent of a 100,000-image source and 1 percent of a one-million-image source. Whether a fixed budget still helps at 100,000 or one million images was not tested.

## F ASSETS AND TERMS OF USE

ImageNet-100-LT is derived from ImageNet, whose access terms restrict the images to non-commercial research and educational use. PASS is released under Creative Commons Attribution 4.0 and DiffusionDB under CC0-1.0. CIFAR-10, CIFAR-100, Stanford Cars, FGVC Aircraft, Oxford Flowers and Oxford-IIIT Pets are used from their official academic releases under their original terms. Stable Diffusion 3 Medium is used under the Stability AI Community License. FLUX.1-dev is used only in the ablation under the FLUX.1-dev Non-Commercial License. We do not redistribute model weights, third-party datasets or generated image corpora.

## G BROADER IMPACT

A method that improves representation quality under biased source data can reduce the need for manual relabeling and can make SSL more useful where balanced data are difficult to collect. The method also relies on large image generators and web-sourced datasets, which carry risks related to license compliance, privacy and unintended memorization or harmful content. Our experiments stay within existing dataset and model release terms, release no new generators or scraped collections, and any deployment would need a review of these risks for its own task.

## H FULL RESULTS

Table 9: Full linear-probe transfer results across all source datasets and comparison methods.

| Source          | Method              | C10                              | C100                             | Cars                             | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|-----------------|---------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| ImageNet-100-LT | SimCLR              | 73.9 $\pm$ 0.5                   | 46.8 $\pm$ 0.4                   | 13.7 $\pm$ 0.3                   | 19.6 $\pm$ 0.2                   | 35.9 $\pm$ 7.2                   | 38.6 $\pm$ 1.5                   | 50.4 $\pm$ 0.5                   | 39.8 $\pm$ 1.5                   |
|                 | TS                  | 72.5 $\pm$ 0.4                   | 44.0 $\pm$ 0.5                   | 12.5 $\pm$ 1.2                   | 15.0 $\pm$ 0.1                   | 35.7 $\pm$ 1.8                   | 40.6 $\pm$ 2.2                   | <b>56.7 <math>\pm</math> 1.2</b> | 39.6 $\pm$ 1.1                   |
|                 | SDCLR               | 41.1 $\pm$ 1.7                   | 18.0 $\pm$ 0.6                   | 1.4 $\pm$ 0.3                    | 2.3 $\pm$ 0.5                    | 10.8 $\pm$ 1.7                   | 10.4 $\pm$ 0.3                   | 21.9 $\pm$ 0.6                   | 15.1 $\pm$ 0.8                   |
|                 | BRIDGE (ours)       | <b>75.8 <math>\pm</math> 0.2</b> | <b>49.3 <math>\pm</math> 0.2</b> | <b>16.4 <math>\pm</math> 1.2</b> | <b>22.4 <math>\pm</math> 0.4</b> | <b>36.8 <math>\pm</math> 1.9</b> | <b>42.9 <math>\pm</math> 1.0</b> | 53.8 $\pm$ 0.1                   | <b>42.5 <math>\pm</math> 0.7</b> |
|                 | BRIDGE+TS (ours)    | 74.2 $\pm$ 0.5                   | 45.6 $\pm$ 0.8                   | 9.4 $\pm$ 0.5                    | 15.0 $\pm$ 0.2                   | 25.5 $\pm$ 2.5                   | 33.1 $\pm$ 2.8                   | 48.3 $\pm$ 0.4                   | 35.9 $\pm$ 1.1                   |
|                 | BRIDGE+SDCLR (ours) | 43.8 $\pm$ 1.7                   | 18.4 $\pm$ 0.7                   | 2.0 $\pm$ 0.3                    | 2.5 $\pm$ 0.8                    | 10.0 $\pm$ 0.8                   | 12.2 $\pm$ 3.1                   | 21.8 $\pm$ 0.5                   | 15.8 $\pm$ 1.1                   |
| CIFAR-10-LT     | SimCLR              | 67.8 $\pm$ 0.2                   | 38.0 $\pm$ 1.0                   | 6.5 $\pm$ 0.4                    | 16.0 $\pm$ 0.9                   | <b>19.2 <math>\pm</math> 1.6</b> | 24.0 $\pm$ 1.3                   | 32.0 $\pm$ 0.7                   | 29.1 $\pm$ 0.9                   |
|                 | TS                  | 68.5 $\pm$ 1.3                   | 36.0 $\pm$ 1.7                   | 4.3 $\pm$ 0.2                    | 10.3 $\pm$ 0.1                   | 13.1 $\pm$ 2.0                   | 19.7 $\pm$ 2.3                   | 31.5 $\pm$ 1.2                   | 26.2 $\pm$ 1.3                   |
|                 | SDCLR               | 39.6 $\pm$ 1.1                   | 15.6 $\pm$ 0.2                   | 1.1 $\pm$ 0.1                    | 2.4 $\pm$ 0.9                    | 12.0 $\pm$ 1.6                   | 9.1 $\pm$ 1.8                    | 16.5 $\pm$ 1.6                   | 13.8 $\pm$ 1.1                   |
|                 | BRIDGE (ours)       | 78.1 $\pm$ 0.3                   | <b>49.8 <math>\pm</math> 0.4</b> | <b>9.2 <math>\pm</math> 0.8</b>  | <b>17.7 <math>\pm</math> 1.0</b> | 18.2 $\pm$ 1.1                   | <b>27.4 <math>\pm</math> 2.7</b> | 39.0 $\pm$ 0.6                   | <b>34.2 <math>\pm</math> 1.0</b> |
|                 | BRIDGE+TS (ours)    | <b>78.6 <math>\pm</math> 0.3</b> | 46.1 $\pm$ 0.3                   | 7.5 $\pm$ 0.9                    | 11.1 $\pm$ 0.7                   | 14.5 $\pm$ 0.6                   | 25.1 $\pm$ 2.7                   | <b>39.6 <math>\pm</math> 0.8</b> | 31.8 $\pm$ 0.9                   |
|                 | BRIDGE+SDCLR (ours) | 40.9 $\pm$ 0.7                   | 17.9 $\pm$ 0.5                   | 1.4 $\pm$ 0.1                    | 2.2 $\pm$ 0.6                    | 11.0 $\pm$ 2.4                   | 10.0 $\pm$ 2.7                   | 21.0 $\pm$ 0.6                   | 14.9 $\pm$ 1.1                   |
| CIFAR-100-LT    | SimCLR              | 59.0 $\pm$ 1.5                   | 31.1 $\pm$ 1.6                   | 4.1 $\pm$ 0.3                    | 11.2 $\pm$ 1.0                   | 12.6 $\pm$ 3.4                   | 20.1 $\pm$ 1.0                   | 28.7 $\pm$ 1.3                   | 23.8 $\pm$ 1.4                   |
|                 | TS                  | 56.3 $\pm$ 0.2                   | 28.0 $\pm$ 0.1                   | 2.8 $\pm$ 0.6                    | 8.5 $\pm$ 0.2                    | 15.5 $\pm$ 2.5                   | 16.2 $\pm$ 3.1                   | 26.2 $\pm$ 1.8                   | 21.9 $\pm$ 1.2                   |
|                 | SDCLR               | 37.1 $\pm$ 0.8                   | 15.7 $\pm$ 0.2                   | 1.9 $\pm$ 0.1                    | 2.4 $\pm$ 0.7                    | 9.4 $\pm$ 1.0                    | 9.9 $\pm$ 1.7                    | 16.2 $\pm$ 0.3                   | 13.2 $\pm$ 0.7                   |
|                 | BRIDGE (ours)       | <b>74.2 <math>\pm</math> 0.5</b> | <b>48.1 <math>\pm</math> 0.4</b> | <b>8.3 <math>\pm</math> 1.0</b>  | <b>18.2 <math>\pm</math> 0.7</b> | <b>24.4 <math>\pm</math> 2.3</b> | 27.6 $\pm$ 0.7                   | 39.8 $\pm$ 1.1                   | <b>34.4 <math>\pm</math> 1.0</b> |
|                 | BRIDGE+TS (ours)    | 73.7 $\pm$ 0.1                   | 47.3 $\pm$ 0.5                   | 7.9 $\pm$ 0.6                    | 14.1 $\pm$ 0.5                   | 18.2 $\pm$ 3.2                   | <b>28.4 <math>\pm</math> 0.4</b> | <b>41.3 <math>\pm</math> 0.3</b> | 33.0 $\pm$ 0.8                   |
|                 | BRIDGE+SDCLR (ours) | 42.0 $\pm$ 1.5                   | 16.9 $\pm$ 0.9                   | 1.3 $\pm$ 0.1                    | 2.2 $\pm$ 0.6                    | 10.6 $\pm$ 0.8                   | 8.3 $\pm$ 0.8                    | 20.2 $\pm$ 0.9                   | 14.5 $\pm$ 0.8                   |
| PASS-10k        | SimCLR              | 72.3 $\pm$ 0.2                   | 45.5 $\pm$ 0.5                   | 11.0 $\pm$ 0.2                   | 17.5 $\pm$ 1.4                   | 33.3 $\pm$ 1.3                   | 33.3 $\pm$ 0.9                   | 44.7 $\pm$ 0.9                   | 36.8 $\pm$ 0.8                   |
|                 | TS                  | 71.5 $\pm$ 0.2                   | 43.5 $\pm$ 0.4                   | 10.2 $\pm$ 1.2                   | 17.2 $\pm$ 0.2                   | 33.1 $\pm$ 2.1                   | 36.4 $\pm$ 1.8                   | 49.4 $\pm$ 0.4                   | 37.3 $\pm$ 0.9                   |
|                 | SDCLR               | 42.2 $\pm$ 1.1                   | 17.7 $\pm$ 0.8                   | 1.8 $\pm$ 0.3                    | 2.2 $\pm$ 0.8                    | 10.6 $\pm$ 1.3                   | 8.2 $\pm$ 1.6                    | 20.8 $\pm$ 0.4                   | 14.8 $\pm$ 0.9                   |
|                 | BRIDGE (ours)       | <b>73.9 <math>\pm</math> 0.3</b> | <b>47.1 <math>\pm</math> 0.4</b> | <b>13.6 <math>\pm</math> 0.7</b> | <b>20.5 <math>\pm</math> 0.7</b> | <b>37.7 <math>\pm</math> 4.0</b> | <b>37.1 <math>\pm</math> 0.7</b> | 49.1 $\pm$ 0.5                   | <b>39.9 <math>\pm</math> 1.0</b> |
|                 | BRIDGE+TS (ours)    | 68.9 $\pm$ 0.6                   | 40.6 $\pm$ 0.7                   | 12.1 $\pm$ 0.4                   | 14.1 $\pm$ 0.5                   | 34.6 $\pm$ 2.9                   | 31.6 $\pm$ 1.5                   | <b>52.4 <math>\pm</math> 0.6</b> | 36.4 $\pm$ 1.0                   |
|                 | BRIDGE+SDCLR (ours) | 44.0 $\pm$ 0.9                   | 18.8 $\pm$ 0.9                   | 1.4 $\pm$ 0.3                    | 1.8 $\pm$ 0.9                    | 12.6 $\pm$ 2.6                   | 12.4 $\pm$ 1.3                   | 21.2 $\pm$ 1.1                   | 16.0 $\pm$ 1.1                   |
| DiffusionDB-10k | SimCLR              | 74.4 $\pm$ 0.5                   | 48.6 $\pm$ 0.2                   | 10.9 $\pm$ 0.7                   | 17.6 $\pm$ 0.8                   | 31.2 $\pm$ 2.4                   | 36.1 $\pm$ 1.3                   | 45.3 $\pm$ 0.5                   | 37.7 $\pm$ 0.9                   |
|                 | TS                  | 72.3 $\pm$ 0.3                   | 44.4 $\pm$ 0.3                   | 10.7 $\pm$ 0.6                   | 15.4 $\pm$ 0.6                   | 33.1 $\pm$ 1.1                   | <b>39.9 <math>\pm</math> 2.7</b> | 48.8 $\pm$ 0.2                   | 37.8 $\pm$ 0.8                   |
|                 | SDCLR               | 42.4 $\pm$ 0.2                   | 18.4 $\pm$ 0.2                   | 1.1 $\pm$ 0.2                    | 2.4 $\pm$ 0.1                    | 13.0 $\pm$ 3.7                   | 8.2 $\pm$ 1.4                    | 21.0 $\pm$ 0.6                   | 15.2 $\pm$ 0.9                   |
|                 | BRIDGE (ours)       | <b>75.9 <math>\pm</math> 0.4</b> | <b>49.9 <math>\pm</math> 0.5</b> | <b>12.8 <math>\pm</math> 1.0</b> | <b>20.0 <math>\pm</math> 0.4</b> | <b>35.7 <math>\pm</math> 1.1</b> | 37.7 $\pm$ 0.8                   | 49.1 $\pm$ 0.5                   | <b>40.1 <math>\pm</math> 0.7</b> |
|                 | BRIDGE+TS (ours)    | 67.8 $\pm$ 0.6                   | 39.1 $\pm$ 0.5                   | 9.4 $\pm$ 1.2                    | 12.4 $\pm$ 0.4                   | 30.3 $\pm$ 3.0                   | 33.7 $\pm$ 1.9                   | <b>50.8 <math>\pm</math> 0.6</b> | 34.8 $\pm$ 1.2                   |
|                 | BRIDGE+SDCLR (ours) | 43.8 $\pm$ 0.9                   | 18.5 $\pm$ 0.3                   | 1.3 $\pm$ 0.4                    | 2.5 $\pm$ 0.0                    | 14.3 $\pm$ 1.6                   | 11.4 $\pm$ 0.6                   | 21.9 $\pm$ 0.6                   | 16.2 $\pm$ 0.6                   |

Table 11:  $k$ NN counterpart of Table 4. Each block changes one component and holds the rest of the loop fixed, the shaded row is the default of that block’s run set, bold marks the best competing setting in a column, and the oracle row is excluded from the bolding.

| Group        | Setting                                 | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|--------------|-----------------------------------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| Generator    | Stable Diffusion 3 (default)            | 68.6 $\pm$ 0.1                   | 37.2 $\pm$ 0.2                   | 6.4 $\pm$ 0.6                   | <b>11.5 <math>\pm</math> 0.2</b> | <b>35.2 <math>\pm</math> 0.3</b> | <b>24.3 <math>\pm</math> 0.6</b> | <b>46.8 <math>\pm</math> 0.7</b> | <b>32.9 <math>\pm</math> 0.4</b> |
|              | FLUX                                    | 65.8 $\pm$ 1.0                   | 35.4 $\pm$ 0.8                   | 5.7 $\pm$ 0.7                   | 9.4 $\pm$ 0.9                    | 27.7 $\pm$ 1.0                   | 18.8 $\pm$ 1.3                   | 38.7 $\pm$ 1.2                   | 28.8 $\pm$ 1.0                   |
|              | Oracle (withheld real images restored)  | 68.8 $\pm$ 0.2                   | 37.7 $\pm$ 0.3                   | 6.6 $\pm$ 0.5                   | 11.4 $\pm$ 0.5                   | 32.7 $\pm$ 0.1                   | 23.6 $\pm$ 0.6                   | 44.4 $\pm$ 0.4                   | 32.2 $\pm$ 0.4                   |
| Selection    | Mode window, lower cutoff p1            | <b>69.2 <math>\pm</math> 0.3</b> | 37.7 $\pm$ 0.2                   | 6.6 $\pm$ 0.2                   | 11.4 $\pm$ 0.2                   | 35.9 $\pm$ 0.2                   | 25.0 $\pm$ 0.5                   | 48.3 $\pm$ 0.3                   | 33.4 $\pm$ 0.3                   |
|              | Mode window, lower cutoff p50           | 68.7 $\pm$ 0.1                   | <b>38.0 <math>\pm</math> 0.2</b> | <b>7.3 <math>\pm</math> 0.2</b> | 11.3 $\pm$ 0.2                   | <b>36.5 <math>\pm</math> 0.1</b> | 24.8 $\pm$ 0.8                   | <b>48.5 <math>\pm</math> 0.4</b> | 33.6 $\pm$ 0.3                   |
|              | Mode window, lower cutoff p75 (default) | 69.0 $\pm$ 0.5                   | 37.8 $\pm$ 0.3                   | 6.7 $\pm$ 0.3                   | <b>11.7 <math>\pm</math> 0.2</b> | 36.5 $\pm$ 0.4                   | <b>25.3 <math>\pm</math> 1.0</b> | 48.1 $\pm$ 0.6                   | <b>33.6 <math>\pm</math> 0.5</b> |
|              | Top tail                                | 65.7 $\pm$ 0.8                   | 34.4 $\pm$ 0.2                   | 4.1 $\pm$ 0.4                   | 9.6 $\pm$ 0.2                    | 23.9 $\pm$ 0.9                   | 18.1 $\pm$ 0.8                   | 35.4 $\pm$ 1.2                   | 27.3 $\pm$ 0.6                   |
|              | Uniform                                 | 66.1 $\pm$ 1.0                   | 35.7 $\pm$ 0.3                   | 4.9 $\pm$ 0.1                   | 10.2 $\pm$ 0.7                   | 25.3 $\pm$ 0.4                   | 19.0 $\pm$ 0.8                   | 37.4 $\pm$ 0.6                   | 28.4 $\pm$ 0.6                   |
| Cycles       | 2 cycles                                | 65.3 $\pm$ 0.4                   | 34.2 $\pm$ 0.5                   | 5.0 $\pm$ 0.7                   | 10.0 $\pm$ 0.2                   | 26.1 $\pm$ 0.8                   | 18.7 $\pm$ 0.6                   | 35.9 $\pm$ 0.3                   | 27.9 $\pm$ 0.5                   |
|              | 5 cycles (default)                      | <b>69.0 <math>\pm</math> 0.5</b> | <b>37.8 <math>\pm</math> 0.3</b> | <b>6.7 <math>\pm</math> 0.3</b> | <b>11.7 <math>\pm</math> 0.2</b> | <b>36.5 <math>\pm</math> 0.4</b> | <b>25.3 <math>\pm</math> 1.0</b> | <b>48.1 <math>\pm</math> 0.6</b> | <b>33.6 <math>\pm</math> 0.5</b> |
|              | 10 cycles                               | 66.4 $\pm$ 0.5                   | 35.4 $\pm$ 0.3                   | 4.8 $\pm$ 0.3                   | 10.4 $\pm$ 0.6                   | 25.2 $\pm$ 0.8                   | 18.6 $\pm$ 0.9                   | 36.0 $\pm$ 0.7                   | 28.1 $\pm$ 0.6                   |
| Architecture | ResNet-18                               | 66.3 $\pm$ 0.2                   | 35.7 $\pm$ 0.3                   | 5.9 $\pm$ 0.2                   | 11.1 $\pm$ 0.4                   | 33.5 $\pm$ 0.1                   | 23.3 $\pm$ 0.5                   | 44.5 $\pm$ 0.7                   | 31.5 $\pm$ 0.4                   |
|              | ResNet-50 (default)                     | 69.0 $\pm$ 0.5                   | 37.8 $\pm$ 0.3                   | <b>6.7 <math>\pm</math> 0.3</b> | <b>11.7 <math>\pm</math> 0.2</b> | 36.5 $\pm$ 0.4                   | <b>25.3 <math>\pm</math> 1.0</b> | <b>48.1 <math>\pm</math> 0.6</b> | <b>33.6 <math>\pm</math> 0.5</b> |
|              | VIT-S                                   | 67.3 $\pm$ 0.4                   | 37.5 $\pm$ 0.3                   | 5.3 $\pm$ 0.2                   | 7.6 $\pm$ 0.1                    | 36.2 $\pm$ 0.7                   | 20.1 $\pm$ 1.0                   | 42.1 $\pm$ 1.0                   | 30.9 $\pm$ 0.5                   |
|              | VIT-B                                   | <b>69.3 <math>\pm</math> 0.8</b> | <b>40.1 <math>\pm</math> 0.8</b> | 6.0 $\pm$ 0.2                   | 8.3 $\pm$ 0.2                    | <b>39.8 <math>\pm</math> 0.8</b> | 21.9 $\pm$ 1.4                   | 45.1 $\pm$ 0.7                   | 32.9 $\pm$ 0.7                   |

Table 10: Full  $k$ NN transfer results across all source datasets and comparison methods.

| Source          | Method              | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|-----------------|---------------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| ImageNet-100-LT | SimCLR              | 68.6 $\pm$ 0.4                   | 37.3 $\pm$ 0.2                   | 6.0 $\pm$ 0.7                   | 11.2 $\pm$ 0.4                   | 32.2 $\pm$ 0.7                   | 23.6 $\pm$ 0.3                   | 43.9 $\pm$ 0.9                   | 31.8 $\pm$ 0.5                   |
|                 | TS                  | 68.4 $\pm$ 0.3                   | 36.5 $\pm$ 0.1                   | 4.7 $\pm$ 0.2                   | 8.3 $\pm$ 0.5                    | 26.6 $\pm$ 1.0                   | 23.6 $\pm$ 1.7                   | <b>52.2 <math>\pm</math> 0.5</b> | 31.5 $\pm$ 0.6                   |
|                 | SDCLR               | 32.5 $\pm$ 0.4                   | 10.9 $\pm$ 0.2                   | 1.0 $\pm$ 0.3                   | 2.1 $\pm$ 0.1                    | 10.7 $\pm$ 0.7                   | 4.7 $\pm$ 0.2                    | 11.2 $\pm$ 0.2                   | 10.4 $\pm$ 0.3                   |
|                 | BRIDGE (ours)       | <b>69.0 <math>\pm</math> 0.5</b> | <b>37.8 <math>\pm</math> 0.3</b> | <b>7.0 <math>\pm</math> 0.2</b> | <b>11.7 <math>\pm</math> 0.2</b> | <b>36.5 <math>\pm</math> 0.4</b> | <b>25.3 <math>\pm</math> 1.0</b> | 48.1 $\pm$ 0.6                   | <b>33.6 <math>\pm</math> 0.4</b> |
|                 | BRIDGE+TS (ours)    | 65.4 $\pm$ 0.2                   | 33.6 $\pm$ 0.3                   | 3.1 $\pm$ 0.1                   | 7.1 $\pm$ 0.3                    | 18.2 $\pm$ 0.2                   | 15.6 $\pm$ 0.7                   | 40.8 $\pm$ 0.5                   | 26.3 $\pm$ 0.3                   |
|                 | BRIDGE+SDCLR (ours) | 36.3 $\pm$ 0.8                   | 13.0 $\pm$ 0.5                   | 1.1 $\pm$ 0.3                   | 2.4 $\pm$ 0.2                    | 9.7 $\pm$ 0.3                    | 5.4 $\pm$ 0.4                    | 13.9 $\pm$ 1.0                   | 11.7 $\pm$ 0.5                   |
| CIFAR-10-LT     | SimCLR              | 66.9 $\pm$ 0.4                   | 33.4 $\pm$ 0.3                   | 3.7 $\pm$ 0.4                   | 9.8 $\pm$ 0.7                    | 16.5 $\pm$ 0.8                   | 13.9 $\pm$ 0.8                   | 24.0 $\pm$ 0.8                   | 24.0 $\pm$ 0.6                   |
|                 | TS                  | 66.3 $\pm$ 0.6                   | 27.8 $\pm$ 0.4                   | 2.4 $\pm$ 0.6                   | 6.4 $\pm$ 0.2                    | 10.5 $\pm$ 0.7                   | 11.1 $\pm$ 0.4                   | 22.0 $\pm$ 0.9                   | 20.9 $\pm$ 0.5                   |
|                 | SDCLR               | 30.1 $\pm$ 0.3                   | 10.0 $\pm$ 0.3                   | 1.7 $\pm$ 0.3                   | 2.2 $\pm$ 0.2                    | 10.1 $\pm$ 0.7                   | 4.4 $\pm$ 0.4                    | 9.8 $\pm$ 0.3                    | 9.7 $\pm$ 0.4                    |
|                 | BRIDGE (ours)       | 74.8 $\pm$ 0.4                   | <b>41.0 <math>\pm</math> 0.4</b> | <b>4.8 <math>\pm</math> 0.5</b> | <b>10.0 <math>\pm</math> 0.2</b> | <b>24.2 <math>\pm</math> 0.6</b> | <b>16.2 <math>\pm</math> 0.7</b> | <b>29.7 <math>\pm</math> 0.3</b> | <b>28.7 <math>\pm</math> 0.4</b> |
|                 | BRIDGE+TS (ours)    | <b>77.6 <math>\pm</math> 0.7</b> | 35.3 $\pm$ 0.5                   | 3.1 $\pm$ 0.3                   | 6.6 $\pm$ 0.4                    | 16.3 $\pm$ 0.6                   | 12.5 $\pm$ 1.0                   | 26.7 $\pm$ 0.1                   | 25.5 $\pm$ 0.5                   |
|                 | BRIDGE+SDCLR (ours) | 34.4 $\pm$ 1.2                   | 11.4 $\pm$ 0.6                   | 1.2 $\pm$ 0.2                   | 2.2 $\pm$ 0.2                    | 8.8 $\pm$ 1.5                    | 5.0 $\pm$ 0.2                    | 11.7 $\pm$ 1.1                   | 10.7 $\pm$ 0.7                   |
| CIFAR-100-LT    | SimCLR              | 59.3 $\pm$ 0.8                   | 30.1 $\pm$ 0.9                   | 3.0 $\pm$ 0.3                   | 7.4 $\pm$ 0.4                    | 13.7 $\pm$ 0.2                   | 10.8 $\pm$ 0.4                   | 21.8 $\pm$ 1.1                   | 20.9 $\pm$ 0.6                   |
|                 | TS                  | 53.1 $\pm$ 0.6                   | 24.6 $\pm$ 0.4                   | 2.4 $\pm$ 0.4                   | 5.5 $\pm$ 0.3                    | 9.3 $\pm$ 0.4                    | 8.3 $\pm$ 0.2                    | 18.7 $\pm$ 1.0                   | 17.4 $\pm$ 0.5                   |
|                 | SDCLR               | 25.7 $\pm$ 1.2                   | 8.5 $\pm$ 0.6                    | 0.9 $\pm$ 0.4                   | 1.9 $\pm$ 0.2                    | 7.3 $\pm$ 0.3                    | 3.9 $\pm$ 0.3                    | 8.3 $\pm$ 1.1                    | 8.1 $\pm$ 0.6                    |
|                 | BRIDGE (ours)       | <b>69.3 <math>\pm</math> 0.5</b> | <b>39.5 <math>\pm</math> 0.1</b> | <b>4.8 <math>\pm</math> 0.6</b> | <b>8.9 <math>\pm</math> 0.1</b>  | <b>21.9 <math>\pm</math> 0.1</b> | <b>15.5 <math>\pm</math> 0.1</b> | <b>29.1 <math>\pm</math> 0.7</b> | <b>27.0 <math>\pm</math> 0.3</b> |
|                 | BRIDGE+TS (ours)    | 64.3 $\pm$ 0.4                   | 36.6 $\pm$ 0.4                   | 2.9 $\pm$ 0.3                   | 6.3 $\pm$ 0.1                    | 16.4 $\pm$ 0.5                   | 12.3 $\pm$ 0.0                   | 28.2 $\pm$ 0.6                   | 23.9 $\pm$ 0.3                   |
|                 | BRIDGE+SDCLR (ours) | 31.9 $\pm$ 0.4                   | 10.4 $\pm$ 0.5                   | 0.8 $\pm$ 0.1                   | 2.1 $\pm$ 0.5                    | 7.5 $\pm$ 1.6                    | 4.5 $\pm$ 0.1                    | 9.1 $\pm$ 0.7                    | 9.5 $\pm$ 0.6                    |
| PASS-10k        | SimCLR              | 67.2 $\pm$ 0.2                   | <b>37.0 <math>\pm</math> 0.3</b> | 5.4 $\pm$ 0.3                   | 10.3 $\pm$ 0.5                   | 30.4 $\pm$ 0.0                   | 16.7 $\pm$ 0.5                   | 38.3 $\pm$ 0.3                   | 29.3 $\pm$ 0.3                   |
|                 | TS                  | 65.2 $\pm$ 0.7                   | 34.6 $\pm$ 0.3                   | 4.4 $\pm$ 0.2                   | 7.9 $\pm$ 0.2                    | 24.5 $\pm$ 0.1                   | 14.4 $\pm$ 0.4                   | 39.3 $\pm$ 0.4                   | 27.2 $\pm$ 0.3                   |
|                 | SDCLR               | 32.4 $\pm$ 1.5                   | 11.5 $\pm$ 0.8                   | 1.3 $\pm$ 0.2                   | 2.0 $\pm$ 0.1                    | 8.6 $\pm$ 0.7                    | 5.2 $\pm$ 0.2                    | 11.1 $\pm$ 0.8                   | 10.3 $\pm$ 0.6                   |
|                 | BRIDGE (ours)       | <b>67.5 <math>\pm</math> 0.3</b> | 36.9 $\pm$ 0.3                   | <b>5.8 <math>\pm</math> 0.5</b> | <b>10.7 <math>\pm</math> 0.4</b> | <b>33.7 <math>\pm</math> 0.2</b> | <b>17.8 <math>\pm</math> 0.6</b> | <b>41.4 <math>\pm</math> 0.8</b> | <b>30.6 <math>\pm</math> 0.4</b> |
|                 | BRIDGE+TS (ours)    | 58.8 $\pm$ 0.5                   | 27.9 $\pm$ 0.5                   | 3.0 $\pm$ 0.2                   | 6.6 $\pm$ 0.1                    | 23.0 $\pm$ 0.3                   | 13.4 $\pm$ 0.4                   | 38.8 $\pm$ 0.8                   | 24.5 $\pm$ 0.4                   |
|                 | BRIDGE+SDCLR (ours) | 37.1 $\pm$ 2.0                   | 13.6 $\pm$ 1.0                   | 1.6 $\pm$ 0.2                   | 2.4 $\pm$ 0.2                    | 10.1 $\pm$ 0.5                   | 6.3 $\pm$ 0.3                    | 13.1 $\pm$ 0.6                   | 12.0 $\pm$ 0.7                   |
| DiffusionDB-10k | SimCLR              | 67.3 $\pm$ 0.1                   | 36.2 $\pm$ 0.4                   | 4.3 $\pm$ 0.2                   | 8.8 $\pm$ 0.1                    | 26.8 $\pm$ 0.4                   | 17.3 $\pm$ 0.1                   | 36.7 $\pm$ 0.8                   | 28.2 $\pm$ 0.3                   |
|                 | TS                  | 64.7 $\pm$ 0.4                   | 33.6 $\pm$ 0.3                   | 3.1 $\pm$ 0.5                   | 7.3 $\pm$ 0.3                    | 23.2 $\pm$ 0.4                   | 17.0 $\pm$ 0.0                   | 37.8 $\pm$ 0.6                   | 26.7 $\pm$ 0.4                   |
|                 | SDCLR               | 32.6 $\pm$ 1.2                   | 10.0 $\pm$ 0.3                   | 1.3 $\pm$ 0.1                   | 2.2 $\pm$ 0.2                    | 6.1 $\pm$ 0.7                    | 4.8 $\pm$ 0.4                    | 10.6 $\pm$ 0.5                   | 9.7 $\pm$ 0.5                    |
|                 | BRIDGE (ours)       | <b>68.3 <math>\pm</math> 0.2</b> | <b>37.4 <math>\pm</math> 0.4</b> | <b>5.4 <math>\pm</math> 0.3</b> | <b>9.8 <math>\pm</math> 0.1</b>  | <b>30.4 <math>\pm</math> 0.2</b> | <b>18.6 <math>\pm</math> 0.2</b> | <b>39.4 <math>\pm</math> 1.0</b> | <b>29.9 <math>\pm</math> 0.3</b> |
|                 | BRIDGE+TS (ours)    | 57.5 $\pm$ 0.8                   | 26.6 $\pm$ 0.3                   | 3.0 $\pm$ 0.4                   | 6.5 $\pm$ 0.1                    | 21.7 $\pm$ 0.8                   | 15.5 $\pm$ 0.4                   | 37.0 $\pm$ 0.4                   | 24.0 $\pm$ 0.5                   |
|                 | BRIDGE+SDCLR (ours) | 35.2 $\pm$ 0.7                   | 11.1 $\pm$ 0.3                   | 1.0 $\pm$ 0.1                   | 2.2 $\pm$ 0.1                    | 8.1 $\pm$ 1.0                    | 5.0 $\pm$ 0.3                    | 11.3 $\pm$ 0.7                   | 10.6 $\pm$ 0.5                   |

Table 12: Augmentation mechanism ablation with full linear-probe results. The oracle row restores the withheld real images.

| Setting                                | C10                              | C100                             | Cars                             | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|----------------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| Stable Diffusion 3 (default)           | <b>75.0 <math>\pm</math> 0.3</b> | <b>48.5 <math>\pm</math> 0.2</b> | <b>16.3 <math>\pm</math> 0.8</b> | <b>20.7 <math>\pm</math> 0.4</b> | 35.7 $\pm$ 0.9                   | <b>41.8 <math>\pm</math> 0.8</b> | <b>52.4 <math>\pm</math> 0.7</b> | <b>41.5 <math>\pm</math> 0.6</b> |
| FLUX                                   | 70.7 $\pm$ 0.4                   | 42.1 $\pm$ 1.4                   | 10.3 $\pm$ 1.3                   | 17.5 $\pm$ 0.9                   | 29.4 $\pm$ 4.6                   | 31.5 $\pm$ 2.9                   | 43.7 $\pm$ 1.4                   | 35.0 $\pm$ 1.9                   |
| Oracle (withheld real images restored) | 74.8 $\pm$ 0.4                   | 47.8 $\pm$ 0.7                   | 13.1 $\pm$ 0.1                   | 20.5 $\pm$ 0.5                   | <b>40.3 <math>\pm</math> 6.1</b> | 38.1 $\pm$ 1.5                   | 50.3 $\pm$ 0.4                   | 40.7 $\pm$ 1.4                   |

Table 13: Augmentation mechanism ablation with full  $k$ NN results. The oracle row restores the withheld real images.

| Setting                                | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|----------------------------------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| Stable Diffusion 3 (default)           | 68.6 $\pm$ 0.1                   | 37.2 $\pm$ 0.2                   | 6.4 $\pm$ 0.6                   | <b>11.5 <math>\pm</math> 0.2</b> | <b>35.2 <math>\pm</math> 0.3</b> | <b>24.3 <math>\pm</math> 0.6</b> | <b>46.8 <math>\pm</math> 0.7</b> | <b>32.9 <math>\pm</math> 0.4</b> |
| FLUX                                   | 65.8 $\pm$ 1.0                   | 35.4 $\pm$ 0.8                   | 5.7 $\pm$ 0.7                   | 9.4 $\pm$ 0.9                    | 27.7 $\pm$ 1.0                   | 18.8 $\pm$ 1.3                   | 38.7 $\pm$ 1.2                   | 28.8 $\pm$ 1.0                   |
| Oracle (withheld real images restored) | <b>68.8 <math>\pm</math> 0.2</b> | <b>37.7 <math>\pm</math> 0.3</b> | <b>6.6 <math>\pm</math> 0.5</b> | 11.4 $\pm$ 0.5                   | 32.7 $\pm$ 0.1                   | 23.6 $\pm$ 0.6                   | 44.4 $\pm$ 0.4                   | 32.2 $\pm$ 0.4                   |

Table 14: Selection rule ablation with full linear-probe results.

| Setting                                  | C10                              | C100                             | Cars                             | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|------------------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| Mode-window q1-99 + diversity            | 75.6 $\pm$ 0.1                   | 49.2 $\pm$ 0.2                   | <b>16.9 <math>\pm</math> 0.7</b> | 21.9 $\pm$ 0.4                   | 34.6 $\pm$ 2.9                   | 43.5 $\pm$ 0.7                   | 52.8 $\pm$ 0.1                   | 42.1 $\pm$ 0.7                   |
| Mode-window q50-99 + diversity           | 75.5 $\pm$ 0.1                   | <b>49.3 <math>\pm</math> 0.4</b> | 15.2 $\pm$ 0.3                   | <b>22.4 <math>\pm</math> 0.6</b> | 35.5 $\pm$ 1.1                   | <b>43.8 <math>\pm</math> 3.1</b> | 53.4 $\pm$ 0.9                   | 42.2 $\pm$ 0.9                   |
| Mode-window q75-99 + diversity (default) | <b>76.0 <math>\pm</math> 0.3</b> | 49.1 $\pm$ 0.2                   | 15.2 $\pm$ 1.2                   | 22.4 $\pm$ 0.4                   | <b>36.8 <math>\pm</math> 1.9</b> | 42.9 $\pm$ 1.0                   | <b>53.8 <math>\pm</math> 0.1</b> | <b>42.3 <math>\pm</math> 0.7</b> |
| Top-tail                                 | 70.8 $\pm$ 0.5                   | 43.1 $\pm$ 1.0                   | 10.0 $\pm$ 1.7                   | 15.9 $\pm$ 1.3                   | 27.7 $\pm$ 1.0                   | 30.5 $\pm$ 1.2                   | 42.2 $\pm$ 1.4                   | 34.3 $\pm$ 1.2                   |
| Uniform                                  | 71.5 $\pm$ 0.2                   | 43.8 $\pm$ 0.7                   | 9.3 $\pm$ 0.9                    | 17.5 $\pm$ 1.2                   | 32.9 $\pm$ 0.4                   | 31.7 $\pm$ 3.9                   | 42.0 $\pm$ 0.9                   | 35.5 $\pm$ 1.2                   |

Table 15: Selection rule ablation with full  $k$ NN results.

| Setting                                  | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|------------------------------------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| Mode-window q1-99 + diversity            | <b>69.2 <math>\pm</math> 0.3</b> | 37.7 $\pm$ 0.2                   | 6.6 $\pm$ 0.2                   | 11.4 $\pm$ 0.2                   | 35.9 $\pm$ 0.2                   | 25.0 $\pm$ 0.5                   | 48.3 $\pm$ 0.3                   | 33.4 $\pm$ 0.3                   |
| Mode-window q50-99 + diversity           | 68.7 $\pm$ 0.1                   | <b>38.0 <math>\pm</math> 0.2</b> | <b>7.3 <math>\pm</math> 0.2</b> | 11.3 $\pm$ 0.2                   | <b>36.5 <math>\pm</math> 0.1</b> | 24.8 $\pm$ 0.8                   | <b>48.5 <math>\pm</math> 0.4</b> | 33.6 $\pm$ 0.3                   |
| Mode-window q75-99 + diversity (default) | 69.0 $\pm$ 0.5                   | 37.8 $\pm$ 0.3                   | 6.7 $\pm$ 0.3                   | <b>11.7 <math>\pm</math> 0.2</b> | 36.5 $\pm$ 0.4                   | <b>25.3 <math>\pm</math> 1.0</b> | 48.1 $\pm$ 0.6                   | <b>33.6 <math>\pm</math> 0.5</b> |
| Top-tail                                 | 65.7 $\pm$ 0.8                   | 34.4 $\pm$ 0.2                   | 4.1 $\pm$ 0.4                   | 9.6 $\pm$ 0.2                    | 23.9 $\pm$ 0.9                   | 18.1 $\pm$ 0.8                   | 35.4 $\pm$ 1.2                   | 27.3 $\pm$ 0.6                   |
| Uniform                                  | 66.1 $\pm$ 1.0                   | 35.7 $\pm$ 0.3                   | 4.9 $\pm$ 0.1                   | 10.2 $\pm$ 0.7                   | 25.3 $\pm$ 0.4                   | 19.0 $\pm$ 0.8                   | 37.4 $\pm$ 0.6                   | 28.4 $\pm$ 0.6                   |

Table 16: Cycle count ablation with full linear-probe results.

| Setting            | C10                              | C100                             | Cars                             | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|--------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| 2 cycles           | 68.9 $\pm$ 0.6                   | 41.6 $\pm$ 1.3                   | 10.6 $\pm$ 1.2                   | 16.4 $\pm$ 0.4                   | 30.7 $\pm$ 4.5                   | 33.2 $\pm$ 0.9                   | 43.3 $\pm$ 0.6                   | 35.0 $\pm$ 1.4                   |
| 5 cycles (default) | <b>76.0 <math>\pm</math> 0.3</b> | <b>49.1 <math>\pm</math> 0.2</b> | <b>15.2 <math>\pm</math> 1.2</b> | <b>22.4 <math>\pm</math> 0.4</b> | <b>36.8 <math>\pm</math> 1.9</b> | <b>42.9 <math>\pm</math> 1.0</b> | <b>53.8 <math>\pm</math> 0.1</b> | <b>42.3 <math>\pm</math> 0.7</b> |
| 10 cycles          | 70.5 $\pm$ 1.2                   | 43.8 $\pm$ 0.6                   | 9.2 $\pm$ 0.9                    | 17.2 $\pm$ 1.6                   | 24.8 $\pm$ 5.2                   | 31.3 $\pm$ 1.7                   | 41.9 $\pm$ 0.6                   | 34.1 $\pm$ 1.7                   |

Table 17: Cycle count ablation with full  $k$ NN results.

| Setting            | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|--------------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| 2 cycles           | 65.3 $\pm$ 0.4                   | 34.2 $\pm$ 0.5                   | 5.0 $\pm$ 0.7                   | 10.0 $\pm$ 0.2                   | 26.1 $\pm$ 0.8                   | 18.7 $\pm$ 0.6                   | 35.9 $\pm$ 0.3                   | 27.9 $\pm$ 0.5                   |
| 5 cycles (default) | <b>69.0 <math>\pm</math> 0.5</b> | <b>37.8 <math>\pm</math> 0.3</b> | <b>6.7 <math>\pm</math> 0.3</b> | <b>11.7 <math>\pm</math> 0.2</b> | <b>36.5 <math>\pm</math> 0.4</b> | <b>25.3 <math>\pm</math> 1.0</b> | <b>48.1 <math>\pm</math> 0.6</b> | <b>33.6 <math>\pm</math> 0.5</b> |
| 10 cycles          | 66.4 $\pm$ 0.5                   | 35.4 $\pm$ 0.3                   | 4.8 $\pm$ 0.3                   | 10.4 $\pm$ 0.6                   | 25.2 $\pm$ 0.8                   | 18.6 $\pm$ 0.9                   | 36.0 $\pm$ 0.7                   | 28.1 $\pm$ 0.6                   |

Table 18: Backbone architecture ablation with full linear-probe results.

| Setting             | C10                              | C100                             | Cars                             | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|---------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| ResNet-18           | 68.8 $\pm$ 0.1                   | 41.7 $\pm$ 0.2                   | 13.5 $\pm$ 0.3                   | 19.7 $\pm$ 0.6                   | 38.7 $\pm$ 3.4                   | 38.4 $\pm$ 2.6                   | 50.2 $\pm$ 0.2                   | 38.7 $\pm$ 1.1                   |
| ResNet-50 (default) | <b>76.0 <math>\pm</math> 0.3</b> | <b>49.1 <math>\pm</math> 0.2</b> | <b>15.2 <math>\pm</math> 1.2</b> | <b>22.4 <math>\pm</math> 0.4</b> | 36.8 $\pm$ 1.9                   | <b>42.9 <math>\pm</math> 1.0</b> | <b>53.8 <math>\pm</math> 0.1</b> | <b>42.3 <math>\pm</math> 0.7</b> |
| ViT-S               | 62.7 $\pm$ 0.4                   | 37.6 $\pm$ 0.6                   | 8.8 $\pm$ 0.5                    | 12.2 $\pm$ 0.2                   | 35.1 $\pm$ 4.2                   | 30.2 $\pm$ 1.6                   | 42.5 $\pm$ 1.2                   | 32.7 $\pm$ 1.3                   |
| ViT-B               | 66.0 $\pm$ 2.3                   | 40.6 $\pm$ 2.3                   | 8.3 $\pm$ 0.5                    | 14.3 $\pm$ 0.8                   | <b>38.8 <math>\pm</math> 2.6</b> | 31.3 $\pm$ 2.1                   | 44.5 $\pm$ 1.7                   | 34.8 $\pm$ 1.8                   |

Table 19: Backbone architecture ablation with full  $k$ NN results.

| Setting             | C10                              | C100                             | Cars                            | Aircraft                         | Flowers                          | Pets                             | IN100-LT                         | Avg.                             |
|---------------------|----------------------------------|----------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|----------------------------------|
| ResNet-18           | 66.3 $\pm$ 0.2                   | 35.7 $\pm$ 0.3                   | 5.9 $\pm$ 0.2                   | 11.1 $\pm$ 0.4                   | 33.5 $\pm$ 0.1                   | 23.3 $\pm$ 0.5                   | 44.5 $\pm$ 0.7                   | 31.5 $\pm$ 0.4                   |
| ResNet-50 (default) | 69.0 $\pm$ 0.5                   | 37.8 $\pm$ 0.3                   | <b>6.7 <math>\pm</math> 0.3</b> | <b>11.7 <math>\pm</math> 0.2</b> | 36.5 $\pm$ 0.4                   | <b>25.3 <math>\pm</math> 1.0</b> | <b>48.1 <math>\pm</math> 0.6</b> | <b>33.6 <math>\pm</math> 0.5</b> |
| ViT-S               | 67.3 $\pm$ 0.4                   | 37.5 $\pm$ 0.3                   | 5.3 $\pm$ 0.2                   | 7.6 $\pm$ 0.1                    | 36.2 $\pm$ 0.7                   | 20.1 $\pm$ 1.0                   | 42.1 $\pm$ 1.0                   | 30.9 $\pm$ 0.5                   |
| ViT-B               | <b>69.3 <math>\pm</math> 0.8</b> | <b>40.1 <math>\pm</math> 0.8</b> | 6.0 $\pm$ 0.2                   | 8.3 $\pm$ 0.2                    | <b>39.8 <math>\pm</math> 0.8</b> | 21.9 $\pm$ 1.4                   | 45.1 $\pm$ 0.7                   | 32.9 $\pm$ 0.7                   |

Table 20: Sensitivity of the acquisition rule on ImageNet-100-LT. Each row changes one component of the acquisition rule, either the candidate-pool multiplier, the upper cutoff, the neighborhood size, the distance or the selection strategy, and holds the rest of the loop fixed under the single-stage protocol of Table 6, so the values are on a different scale from the five-cycle tables. Mean  $\pm$  standard deviation over three seeds. Cosine scoring rescales the normalized  $L_2$  score by a positive constant and therefore selects the identical anchor set, so the gap between the two rows is run-to-run spread. Linear-probe accuracy. The last column is the mean of the seven task means.

| Method                                | Cars            | Aircraft         | Flowers          | Pets             | CIFAR-10         | CIFAR-100        | IN100-LT         | Avg.  |
|---------------------------------------|-----------------|------------------|------------------|------------------|------------------|------------------|------------------|-------|
| Candidate pool $\alpha = 1$           | 6.46 $\pm$ 0.57 | 11.37 $\pm$ 1.59 | 23.53 $\pm$ 3.00 | 21.92 $\pm$ 2.00 | 61.94 $\pm$ 0.94 | 32.73 $\pm$ 0.21 | 33.92 $\pm$ 1.68 | 27.41 |
| Candidate pool $\alpha = 2$           | 6.77 $\pm$ 0.63 | 11.59 $\pm$ 0.40 | 21.79 $\pm$ 3.22 | 23.19 $\pm$ 3.93 | 62.90 $\pm$ 1.06 | 34.48 $\pm$ 0.11 | 35.69 $\pm$ 0.87 | 28.06 |
| Candidate pool $\alpha = 4$ (default) | 6.31 $\pm$ 0.10 | 11.17 $\pm$ 3.78 | 23.53 $\pm$ 3.27 | 24.18 $\pm$ 2.16 | 62.11 $\pm$ 0.19 | 33.39 $\pm$ 0.64 | 34.18 $\pm$ 0.47 | 27.84 |
| Candidate pool $\alpha = 8$           | 6.41 $\pm$ 0.28 | 11.16 $\pm$ 0.78 | 23.09 $\pm$ 4.35 | 24.82 $\pm$ 3.29 | 62.60 $\pm$ 1.79 | 33.15 $\pm$ 1.01 | 35.40 $\pm$ 0.87 | 28.09 |
| Upper cutoff $q = 0.95$               | 6.79 $\pm$ 0.65 | 11.44 $\pm$ 0.85 | 19.83 $\pm$ 2.47 | 25.27 $\pm$ 2.32 | 62.57 $\pm$ 0.87 | 33.83 $\pm$ 0.53 | 35.13 $\pm$ 1.03 | 27.84 |
| Upper cutoff $q = 0.99$ (default)     | 6.59 $\pm$ 0.17 | 11.76 $\pm$ 1.54 | 20.26 $\pm$ 1.73 | 24.00 $\pm$ 2.20 | 63.11 $\pm$ 0.75 | 33.43 $\pm$ 1.02 | 35.85 $\pm$ 1.15 | 27.86 |
| Upper cutoff $q = 0.995$              | 6.78 $\pm$ 0.66 | 10.40 $\pm$ 2.49 | 22.88 $\pm$ 4.58 | 22.55 $\pm$ 3.09 | 62.76 $\pm$ 0.50 | 33.87 $\pm$ 0.38 | 35.32 $\pm$ 0.64 | 27.79 |
| Upper cutoff $q = 1.0$                | 6.31 $\pm$ 0.45 | 10.94 $\pm$ 1.34 | 22.44 $\pm$ 1.89 | 23.37 $\pm$ 4.00 | 62.85 $\pm$ 1.04 | 33.91 $\pm$ 0.87 | 35.51 $\pm$ 1.59 | 27.90 |
| Neighborhood $k = 10$                 | 6.26 $\pm$ 0.29 | 10.44 $\pm$ 1.43 | 22.22 $\pm$ 0.65 | 24.00 $\pm$ 1.23 | 62.66 $\pm$ 0.60 | 33.44 $\pm$ 0.58 | 35.14 $\pm$ 1.21 | 27.74 |
| Neighborhood $k = 25$                 | 6.44 $\pm$ 0.76 | 11.13 $\pm$ 1.19 | 19.39 $\pm$ 3.60 | 23.37 $\pm$ 1.18 | 63.18 $\pm$ 0.56 | 33.40 $\pm$ 1.21 | 33.49 $\pm$ 0.43 | 27.20 |
| Neighborhood $k = 50$                 | 6.34 $\pm$ 0.56 | 10.01 $\pm$ 1.32 | 20.48 $\pm$ 1.00 | 22.64 $\pm$ 3.24 | 61.28 $\pm$ 2.19 | 32.64 $\pm$ 0.29 | 34.23 $\pm$ 0.77 | 26.80 |
| Neighborhood $k = 200$                | 6.50 $\pm$ 0.46 | 10.58 $\pm$ 2.83 | 20.26 $\pm$ 1.13 | 22.46 $\pm$ 1.23 | 61.24 $\pm$ 1.08 | 32.87 $\pm$ 0.76 | 34.42 $\pm$ 1.23 | 26.90 |
| Distance: normalized $L_2$ (default)  | 6.59 $\pm$ 0.17 | 11.79 $\pm$ 1.34 | 20.48 $\pm$ 2.64 | 24.46 $\pm$ 2.68 | 61.69 $\pm$ 0.67 | 33.88 $\pm$ 1.05 | 35.24 $\pm$ 1.10 | 27.73 |
| Distance: cosine (null control)       | 6.55 $\pm$ 0.15 | 10.50 $\pm$ 1.10 | 20.04 $\pm$ 2.64 | 23.91 $\pm$ 1.70 | 62.22 $\pm$ 1.29 | 33.23 $\pm$ 0.81 | 34.98 $\pm$ 0.51 | 27.35 |
| Distance: raw $L_2$                   | 6.92 $\pm$ 0.78 | 10.45 $\pm$ 3.00 | 18.95 $\pm$ 5.23 | 25.72 $\pm$ 2.31 | 62.81 $\pm$ 0.80 | 33.20 $\pm$ 0.89 | 35.08 $\pm$ 2.48 | 27.59 |
| Strategy: mode window (default)       | 6.55 $\pm$ 0.15 | 10.18 $\pm$ 1.30 | 22.44 $\pm$ 4.59 | 23.28 $\pm$ 2.04 | 60.95 $\pm$ 3.01 | 33.05 $\pm$ 0.21 | 34.53 $\pm$ 0.92 | 27.28 |
| Strategy: sparse band with FPS        | 6.60 $\pm$ 0.62 | 11.21 $\pm$ 1.26 | 21.13 $\pm$ 2.64 | 22.46 $\pm$ 0.16 | 63.34 $\pm$ 0.27 | 34.29 $\pm$ 0.30 | 34.30 $\pm$ 1.26 | 27.62 |
| Strategy: sparse band, random         | 6.46 $\pm$ 0.48 | 11.06 $\pm$ 0.70 | 21.35 $\pm$ 4.82 | 22.92 $\pm$ 0.87 | 62.26 $\pm$ 0.89 | 32.33 $\pm$ 0.96 | 34.46 $\pm$ 1.08 | 27.26 |
| Strategy: inverse-cluster             | 6.44 $\pm$ 0.75 | 10.66 $\pm$ 1.06 | 19.17 $\pm$ 2.47 | 21.65 $\pm$ 2.28 | 62.70 $\pm$ 0.45 | 33.50 $\pm$ 0.95 | 34.69 $\pm$ 1.10 | 26.97 |
| Strategy: densest window              | 6.64 $\pm$ 0.46 | 11.90 $\pm$ 0.49 | 20.04 $\pm$ 1.00 | 24.28 $\pm$ 1.64 | 63.50 $\pm$ 0.60 | 33.80 $\pm$ 0.93 | 35.20 $\pm$ 0.80 | 27.91 |
| Strategy: top tail, no diversity      | 6.39 $\pm$ 0.43 | 10.48 $\pm$ 1.03 | 19.83 $\pm$ 2.00 | 23.82 $\pm$ 1.50 | 62.86 $\pm$ 0.92 | 33.14 $\pm$ 1.02 | 34.66 $\pm$ 0.41 | 27.31 |

Table 21: Sensitivity of the acquisition rule on ImageNet-100-LT. Each row changes one component of the acquisition rule, either the candidate-pool multiplier, the upper cutoff, the neighborhood size, the distance or the selection strategy, and holds the rest of the loop fixed under the single-stage protocol of Table 6, so the values are on a different scale from the five-cycle tables. Mean  $\pm$  standard deviation over three seeds. Cosine scoring rescales the normalized  $L_2$  score by a positive constant and therefore selects the identical anchor set, so the gap between the two rows is run-to-run spread.  $k$ NN accuracy. The last column is the mean of the seven task means.

| Method                                | Cars            | Aircraft        | Flowers          | Pets             | CIFAR-10         | CIFAR-100        | IN100-LT         | Avg.  |
|---------------------------------------|-----------------|-----------------|------------------|------------------|------------------|------------------|------------------|-------|
| Candidate pool $\alpha = 1$           | 5.02 $\pm$ 0.14 | 8.62 $\pm$ 0.26 | 19.22 $\pm$ 0.71 | 14.92 $\pm$ 0.99 | 62.40 $\pm$ 0.12 | 31.47 $\pm$ 0.25 | 33.47 $\pm$ 1.05 | 25.02 |
| Candidate pool $\alpha = 2$           | 5.08 $\pm$ 0.15 | 8.36 $\pm$ 0.22 | 19.07 $\pm$ 0.42 | 15.17 $\pm$ 0.27 | 62.81 $\pm$ 0.58 | 31.72 $\pm$ 0.27 | 33.19 $\pm$ 0.61 | 25.06 |
| Candidate pool $\alpha = 4$ (default) | 4.92 $\pm$ 0.15 | 8.48 $\pm$ 0.72 | 19.04 $\pm$ 0.03 | 14.60 $\pm$ 0.38 | 62.35 $\pm$ 0.54 | 31.84 $\pm$ 0.37 | 32.59 $\pm$ 0.49 | 24.83 |
| Candidate pool $\alpha = 8$           | 4.85 $\pm$ 0.17 | 8.53 $\pm$ 0.41 | 19.17 $\pm$ 0.90 | 14.82 $\pm$ 0.62 | 62.15 $\pm$ 0.07 | 31.56 $\pm$ 0.23 | 33.58 $\pm$ 1.03 | 24.95 |
| Upper cutoff $q = 0.95$               | 4.99 $\pm$ 0.11 | 8.57 $\pm$ 0.51 | 19.09 $\pm$ 0.51 | 14.68 $\pm$ 0.28 | 62.84 $\pm$ 0.45 | 31.88 $\pm$ 0.04 | 32.56 $\pm$ 0.45 | 24.94 |
| Upper cutoff $q = 0.99$ (default)     | 4.95 $\pm$ 0.17 | 8.53 $\pm$ 0.69 | 19.40 $\pm$ 0.45 | 14.85 $\pm$ 0.49 | 62.29 $\pm$ 0.50 | 31.87 $\pm$ 0.16 | 32.44 $\pm$ 0.48 | 24.90 |
| Upper cutoff $q = 0.995$              | 5.01 $\pm$ 0.07 | 8.34 $\pm$ 0.74 | 19.58 $\pm$ 0.83 | 14.87 $\pm$ 0.48 | 63.05 $\pm$ 0.47 | 32.06 $\pm$ 0.11 | 33.01 $\pm$ 0.34 | 25.13 |
| Upper cutoff $q = 1.0$                | 4.99 $\pm$ 0.08 | 8.15 $\pm$ 0.61 | 18.52 $\pm$ 0.98 | 14.45 $\pm$ 0.19 | 62.84 $\pm$ 0.32 | 31.69 $\pm$ 0.30 | 32.99 $\pm$ 0.83 | 24.80 |
| Neighborhood $k = 10$                 | 5.04 $\pm$ 0.09 | 8.99 $\pm$ 1.21 | 18.86 $\pm$ 0.35 | 14.90 $\pm$ 0.57 | 62.73 $\pm$ 0.39 | 31.69 $\pm$ 0.17 | 33.14 $\pm$ 0.40 | 25.05 |
| Neighborhood $k = 25$                 | 4.97 $\pm$ 0.14 | 8.59 $\pm$ 0.21 | 19.05 $\pm$ 0.29 | 15.04 $\pm$ 0.85 | 62.68 $\pm$ 0.66 | 31.82 $\pm$ 0.16 | 33.56 $\pm$ 0.97 | 25.10 |
| Neighborhood $k = 50$                 | 5.06 $\pm$ 0.17 | 8.62 $\pm$ 0.77 | 19.26 $\pm$ 0.67 | 14.92 $\pm$ 0.03 | 62.63 $\pm$ 0.40 | 31.71 $\pm$ 0.59 | 31.88 $\pm$ 0.57 | 24.87 |
| Neighborhood $k = 200$                | 5.17 $\pm$ 0.15 | 8.65 $\pm$ 0.59 | 19.22 $\pm$ 1.02 | 15.17 $\pm$ 0.10 | 62.80 $\pm$ 0.31 | 31.90 $\pm$ 0.13 | 33.91 $\pm$ 0.54 | 25.26 |
| Distance: normalized $L_2$ (default)  | 4.97 $\pm$ 0.18 | 8.55 $\pm$ 0.67 | 19.42 $\pm$ 0.49 | 14.85 $\pm$ 0.50 | 62.27 $\pm$ 0.52 | 31.86 $\pm$ 0.16 | 32.60 $\pm$ 0.96 | 24.93 |
| Distance: cosine (null control)       | 4.90 $\pm$ 0.13 | 8.52 $\pm$ 0.69 | 19.35 $\pm$ 0.54 | 14.75 $\pm$ 0.44 | 62.31 $\pm$ 0.53 | 31.78 $\pm$ 0.28 | 32.44 $\pm$ 0.83 | 24.86 |
| Distance: raw $L_2$                   | 4.85 $\pm$ 0.10 | 7.83 $\pm$ 0.39 | 19.13 $\pm$ 0.88 | 14.63 $\pm$ 0.34 | 62.81 $\pm$ 0.52 | 31.71 $\pm$ 0.04 | 32.28 $\pm$ 0.65 | 24.75 |
| Strategy: mode window (default)       | 4.90 $\pm$ 0.12 | 8.52 $\pm$ 0.69 | 19.34 $\pm$ 0.54 | 14.75 $\pm$ 0.45 | 62.31 $\pm$ 0.52 | 31.78 $\pm$ 0.28 | 32.73 $\pm$ 0.75 | 24.90 |
| Strategy: sparse band with FPS        | 5.19 $\pm$ 0.04 | 8.12 $\pm$ 0.39 | 19.53 $\pm$ 0.49 | 14.95 $\pm$ 0.75 | 62.43 $\pm$ 0.17 | 31.62 $\pm$ 0.32 | 33.38 $\pm$ 1.46 | 25.03 |
| Strategy: sparse band, random         | 5.00 $\pm$ 0.13 | 8.70 $\pm$ 0.06 | 19.36 $\pm$ 0.88 | 14.95 $\pm$ 0.61 | 62.63 $\pm$ 0.19 | 31.64 $\pm$ 0.62 | 33.04 $\pm$ 0.82 | 25.05 |
| Strategy: inverse-cluster             | 5.05 $\pm$ 0.02 | 8.83 $\pm$ 1.05 | 19.40 $\pm$ 0.81 | 15.06 $\pm$ 0.46 | 62.60 $\pm$ 0.50 | 31.99 $\pm$ 0.37 | 33.75 $\pm$ 0.44 | 25.24 |
| Strategy: densest window              | 4.88 $\pm$ 0.10 | 8.44 $\pm$ 0.54 | 19.34 $\pm$ 0.59 | 14.95 $\pm$ 0.53 | 62.46 $\pm$ 0.12 | 31.62 $\pm$ 0.30 | 33.08 $\pm$ 0.64 | 24.97 |
| Strategy: top tail, no diversity      | 4.90 $\pm$ 0.12 | 8.38 $\pm$ 0.29 | 19.11 $\pm$ 0.64 | 15.22 $\pm$ 0.03 | 62.05 $\pm$ 0.77 | 31.77 $\pm$ 0.13 | 32.41 $\pm$ 0.25 | 24.83 |

Table 22: Repair utility by sparsity band on ImageNet-100-LT. Each row spends the identical acquisition budget on a different quantile range of the  $k$ NN score distribution, holding every other part of the loop fixed. Mean  $\pm$  standard deviation over three seeds. Linear-probe accuracy. The last column is the mean of the seven task means.

| Method                 | Cars            | Aircraft         | Flowers          | Pets             | CIFAR-10         | CIFAR-100        | IN100-LT         | Avg.  |
|------------------------|-----------------|------------------|------------------|------------------|------------------|------------------|------------------|-------|
| q0–25 (densest)        | 6.73 $\pm$ 0.36 | 9.63 $\pm$ 0.26  | 24.84 $\pm$ 5.10 | 23.37 $\pm$ 2.72 | 62.97 $\pm$ 0.82 | 33.96 $\pm$ 0.49 | 34.44 $\pm$ 0.62 | 27.99 |
| q25–50                 | 6.80 $\pm$ 0.39 | 10.28 $\pm$ 1.37 | 23.75 $\pm$ 2.10 | 22.74 $\pm$ 1.57 | 62.76 $\pm$ 0.64 | 33.20 $\pm$ 1.12 | 35.61 $\pm$ 0.23 | 27.88 |
| q50–75                 | 6.79 $\pm$ 0.63 | 10.24 $\pm$ 1.94 | 20.04 $\pm$ 4.64 | 23.55 $\pm$ 2.58 | 62.24 $\pm$ 0.37 | 32.80 $\pm$ 1.66 | 34.51 $\pm$ 1.27 | 27.17 |
| q75–85                 | 6.12 $\pm$ 0.33 | 9.74 $\pm$ 1.00  | 22.88 $\pm$ 4.58 | 23.46 $\pm$ 1.64 | 60.97 $\pm$ 1.05 | 31.90 $\pm$ 0.36 | 33.55 $\pm$ 0.46 | 26.95 |
| q85–90                 | 6.62 $\pm$ 0.39 | 9.95 $\pm$ 0.73  | 23.75 $\pm$ 3.09 | 25.45 $\pm$ 1.77 | 63.38 $\pm$ 0.78 | 33.47 $\pm$ 0.93 | 35.24 $\pm$ 0.69 | 28.27 |
| q90–95                 | 6.55 $\pm$ 0.55 | 10.46 $\pm$ 1.87 | 22.66 $\pm$ 3.94 | 25.27 $\pm$ 0.72 | 62.30 $\pm$ 1.15 | 33.49 $\pm$ 0.87 | 34.40 $\pm$ 0.62 | 27.88 |
| q95–97                 | 6.58 $\pm$ 0.20 | 9.55 $\pm$ 1.62  | 23.09 $\pm$ 1.00 | 23.55 $\pm$ 0.68 | 61.77 $\pm$ 2.13 | 32.75 $\pm$ 2.27 | 34.99 $\pm$ 1.19 | 27.47 |
| q97–99                 | 6.80 $\pm$ 0.17 | 11.38 $\pm$ 2.82 | 20.26 $\pm$ 3.98 | 26.27 $\pm$ 3.24 | 62.73 $\pm$ 1.49 | 33.13 $\pm$ 1.97 | 35.31 $\pm$ 1.55 | 27.98 |
| q99–100 (extreme tail) | 6.55 $\pm$ 0.28 | 10.78 $\pm$ 1.23 | 21.79 $\pm$ 0.75 | 23.82 $\pm$ 1.37 | 63.54 $\pm$ 0.92 | 33.65 $\pm$ 0.55 | 35.57 $\pm$ 0.91 | 27.96 |

Table 23: Repair utility by sparsity band on ImageNet-100-LT. Each row spends the identical acquisition budget on a different quantile range of the  $k$ NN score distribution, holding every other part of the loop fixed. Mean  $\pm$  standard deviation over three seeds.  $k$ NN accuracy. The last column is the mean of the seven task means.

| Method                 | Cars            | Aircraft        | Flowers          | Pets             | CIFAR-10         | CIFAR-100        | IN100-LT         | Avg.  |
|------------------------|-----------------|-----------------|------------------|------------------|------------------|------------------|------------------|-------|
| q0–25 (densest)        | 4.97 $\pm$ 0.07 | 8.33 $\pm$ 0.72 | 20.00 $\pm$ 0.65 | 14.48 $\pm$ 0.27 | 62.04 $\pm$ 0.84 | 31.45 $\pm$ 0.05 | 31.88 $\pm$ 0.91 | 24.74 |
| q25–50                 | 5.16 $\pm$ 0.14 | 8.52 $\pm$ 0.37 | 19.58 $\pm$ 0.49 | 14.82 $\pm$ 0.73 | 62.95 $\pm$ 0.79 | 31.54 $\pm$ 0.09 | 32.88 $\pm$ 0.78 | 25.06 |
| q50–75                 | 5.06 $\pm$ 0.08 | 8.51 $\pm$ 0.44 | 19.63 $\pm$ 0.75 | 14.46 $\pm$ 0.81 | 62.41 $\pm$ 0.61 | 31.47 $\pm$ 0.56 | 32.61 $\pm$ 0.55 | 24.88 |
| q75–85                 | 4.97 $\pm$ 0.08 | 8.78 $\pm$ 0.45 | 19.54 $\pm$ 0.14 | 15.39 $\pm$ 0.20 | 62.52 $\pm$ 0.56 | 31.31 $\pm$ 0.43 | 32.23 $\pm$ 1.17 | 24.96 |
| q85–90                 | 5.04 $\pm$ 0.09 | 9.04 $\pm$ 0.48 | 19.61 $\pm$ 0.46 | 14.75 $\pm$ 0.65 | 62.57 $\pm$ 0.27 | 31.74 $\pm$ 0.34 | 32.82 $\pm$ 0.96 | 25.08 |
| q90–95                 | 4.92 $\pm$ 0.21 | 8.67 $\pm$ 0.75 | 19.60 $\pm$ 0.53 | 14.96 $\pm$ 0.46 | 62.40 $\pm$ 0.59 | 31.56 $\pm$ 0.54 | 32.48 $\pm$ 0.56 | 24.94 |
| q95–97                 | 4.88 $\pm$ 0.35 | 8.51 $\pm$ 1.09 | 19.40 $\pm$ 0.72 | 14.97 $\pm$ 0.49 | 62.32 $\pm$ 0.94 | 31.55 $\pm$ 0.43 | 31.63 $\pm$ 0.32 | 24.75 |
| q97–99                 | 5.02 $\pm$ 0.08 | 8.59 $\pm$ 0.43 | 19.63 $\pm$ 0.67 | 14.89 $\pm$ 0.56 | 62.64 $\pm$ 0.56 | 31.75 $\pm$ 0.61 | 32.59 $\pm$ 0.96 | 25.02 |
| q99–100 (extreme tail) | 5.17 $\pm$ 0.16 | 8.16 $\pm$ 0.23 | 19.10 $\pm$ 0.60 | 15.00 $\pm$ 0.53 | 62.81 $\pm$ 0.26 | 32.12 $\pm$ 0.43 | 32.15 $\pm$ 1.40 | 24.93 |

## I ALGORITHM

### Algorithm 1 BRIDGE training loop with mode-window selection

```

1: Initialize source set  $D^{(0)}$ , encoder  $f_{\theta^{(0)}}$ , frozen generator  $g_\phi$ , cycle lengths  $\{E_c\}_{c=0}^{C-1}$ 
2: for  $c = 0, \dots, C - 1$  do
3:   Continue training  $f_{\theta^{(c)}}$  on  $D^{(c)}$  for  $E_c$  epochs with a fresh optimizer schedule
4:   if  $c < C - 1$  then
5:     Compute normalized embeddings  $Z^{(c)} = \{f_{\theta^{(c)}}(x_i) \mid x_i \in D^{(c)}\}$ 
6:     Compute mean  $k$ NN scores  $\{d_i^{(c)}\}$  from  $Z^{(c)}$ 
7:     Restrict to the 75th to 99th percentile band and locate the modal bin center  $\bar{d}^{(c)}$ 
8:     Form the candidate pool  $P^{(c)}$  of size  $\alpha B$  whose scores are closest to  $\bar{d}^{(c)}$ 
9:     Select  $S^{(c)} \subseteq P^{(c)}$  with  $|S^{(c)}| = B$  by greedy farthest-point sampling in embedding space
10:    Form  $\tilde{D}^{(c)} = \{g_\phi(x_i, \epsilon_j) \mid i \in S^{(c)}, j \in [G]\}$  and set  $D^{(c+1)} = D^{(c)} \cup \tilde{D}^{(c)}$ 
11:   end if
12: end for
13: Return the final encoder  $f_{\theta^{(C-1)}}$ 

```

Table 24: Paired full-schedule ViT-S transfer on ImageNet-100-LT, with every downstream dataset reported separately. Both arms of each pair start from the same source checkpoint and receive the same number of post-branch optimizer updates. Four objectives span contrastive (SimCLR, MoCo v3), self-distillation (DINO) and masked-reconstruction (MAE) pretraining. The first three run five cycles of 100 epochs and MAE five of 80. Mean  $\pm$  standard deviation over three seeds. Linear-probe accuracy.

| Method                | Cars            | Aircraft         | Flowers          | Pets             | CIFAR-10         | CIFAR-100        | IN100-LT         |
|-----------------------|-----------------|------------------|------------------|------------------|------------------|------------------|------------------|
| SimCLR, SSL baseline  | 7.15 $\pm$ 0.19 | 10.92 $\pm$ 0.40 | 32.90 $\pm$ 2.72 | 25.00 $\pm$ 1.25 | 60.46 $\pm$ 0.46 | 34.96 $\pm$ 0.20 | 41.95 $\pm$ 1.43 |
| SimCLR, with BRIDGE   | 8.18 $\pm$ 0.42 | 11.82 $\pm$ 0.43 | 36.60 $\pm$ 2.85 | 29.80 $\pm$ 2.98 | 60.85 $\pm$ 0.35 | 36.33 $\pm$ 0.14 | 43.93 $\pm$ 1.09 |
| MoCo v3, SSL baseline | 7.57 $\pm$ 0.23 | 11.05 $\pm$ 1.31 | 30.28 $\pm$ 1.36 | 29.26 $\pm$ 3.13 | 59.96 $\pm$ 0.39 | 34.47 $\pm$ 0.01 | 44.79 $\pm$ 1.17 |
| MoCo v3, with BRIDGE  | 9.55 $\pm$ 0.34 | 13.42 $\pm$ 0.98 | 36.82 $\pm$ 6.07 | 31.25 $\pm$ 4.00 | 62.36 $\pm$ 0.46 | 36.30 $\pm$ 0.13 | 47.55 $\pm$ 0.86 |
| DINO, SSL baseline    | 8.10 $\pm$ 0.25 | 11.40 $\pm$ 1.03 | 38.13 $\pm$ 5.24 | 27.54 $\pm$ 2.34 | 58.93 $\pm$ 0.94 | 33.56 $\pm$ 1.75 | 44.63 $\pm$ 1.47 |
| DINO, with BRIDGE     | 9.49 $\pm$ 0.88 | 13.28 $\pm$ 0.58 | 37.25 $\pm$ 2.61 | 34.60 $\pm$ 3.62 | 61.30 $\pm$ 0.07 | 35.23 $\pm$ 0.31 | 46.78 $\pm$ 0.77 |
| MAE, SSL baseline     | 1.54 $\pm$ 0.19 | 2.44 $\pm$ 1.00  | 10.46 $\pm$ 5.88 | 10.05 $\pm$ 3.03 | 37.69 $\pm$ 3.73 | 13.20 $\pm$ 1.94 | 19.93 $\pm$ 2.13 |
| MAE, with BRIDGE      | 1.56 $\pm$ 0.43 | 3.18 $\pm$ 0.70  | 8.50 $\pm$ 2.26  | 8.61 $\pm$ 2.20  | 41.46 $\pm$ 2.20 | 15.00 $\pm$ 2.40 | 22.19 $\pm$ 2.35 |

Table 25: Paired full-schedule ViT-S transfer on ImageNet-100-LT, with every downstream dataset reported separately. Both arms of each pair start from the same source checkpoint and receive the same number of post-branch optimizer updates. Four objectives span contrastive (SimCLR, MoCo v3), self-distillation (DINO) and masked-reconstruction (MAE) pretraining. The first three run five cycles of 100 epochs and MAE five of 80. Mean  $\pm$  standard deviation over three seeds.  $k$ NN accuracy.

| Method                | Cars            | Aircraft        | Flowers          | Pets             | CIFAR-10         | CIFAR-100        | IN100-LT         |
|-----------------------|-----------------|-----------------|------------------|------------------|------------------|------------------|------------------|
| SimCLR, SSL baseline  | 6.26 $\pm$ 0.17 | 7.09 $\pm$ 0.17 | 36.53 $\pm$ 0.89 | 19.16 $\pm$ 0.64 | 66.05 $\pm$ 0.48 | 36.72 $\pm$ 0.36 | 41.21 $\pm$ 0.72 |
| SimCLR, with BRIDGE   | 6.52 $\pm$ 0.22 | 7.56 $\pm$ 0.23 | 38.29 $\pm$ 0.83 | 19.93 $\pm$ 0.53 | 65.70 $\pm$ 0.85 | 36.46 $\pm$ 0.33 | 42.91 $\pm$ 0.78 |
| MoCo v3, SSL baseline | 6.65 $\pm$ 0.38 | 7.71 $\pm$ 0.16 | 38.91 $\pm$ 1.03 | 23.56 $\pm$ 1.12 | 69.29 $\pm$ 0.25 | 39.98 $\pm$ 0.44 | 45.14 $\pm$ 0.46 |
| MoCo v3, with BRIDGE  | 6.74 $\pm$ 0.26 | 8.21 $\pm$ 0.18 | 40.51 $\pm$ 1.45 | 24.91 $\pm$ 0.45 | 65.41 $\pm$ 0.21 | 36.24 $\pm$ 0.46 | 47.71 $\pm$ 1.53 |
| DINO, SSL baseline    | 6.05 $\pm$ 0.27 | 7.73 $\pm$ 0.52 | 38.52 $\pm$ 1.19 | 19.26 $\pm$ 0.16 | 62.69 $\pm$ 2.42 | 33.98 $\pm$ 2.80 | 45.28 $\pm$ 0.98 |
| DINO, with BRIDGE     | 6.39 $\pm$ 0.22 | 8.16 $\pm$ 0.26 | 40.72 $\pm$ 0.31 | 21.79 $\pm$ 1.15 | 64.38 $\pm$ 3.51 | 35.59 $\pm$ 2.51 | 47.14 $\pm$ 0.29 |
| MAE, SSL baseline     | 1.00 $\pm$ 0.14 | 2.24 $\pm$ 0.29 | 5.30 $\pm$ 0.52  | 4.58 $\pm$ 0.17  | 31.34 $\pm$ 1.65 | 10.96 $\pm$ 0.44 | 11.75 $\pm$ 0.38 |
| MAE, with BRIDGE      | 1.38 $\pm$ 0.37 | 2.50 $\pm$ 0.70 | 7.54 $\pm$ 2.13  | 5.35 $\pm$ 0.76  | 35.58 $\pm$ 7.14 | 13.44 $\pm$ 3.46 | 15.46 $\pm$ 2.68 |